

#### Red Hat OpenShift Container Platform

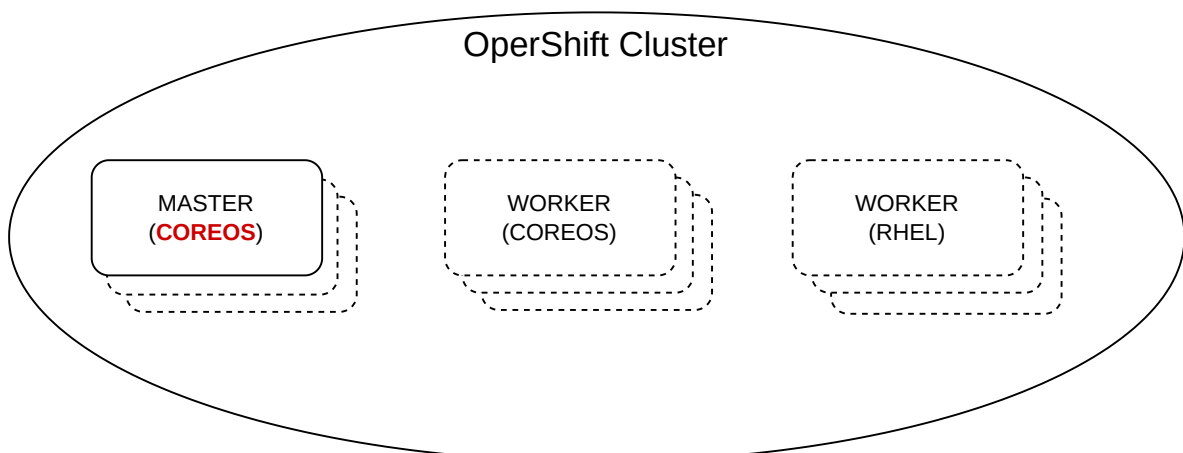
- Public/private DC.
- Bare metal and multiple cloud and virtualization providers.
- Full control by customer.

#### Red Hat OpenShift Dedicated



- Managed cluster in public cloud.
- RH manages the cluster.
- Customer manages updates and add-on services.

#### Red Hat OpenShift Online

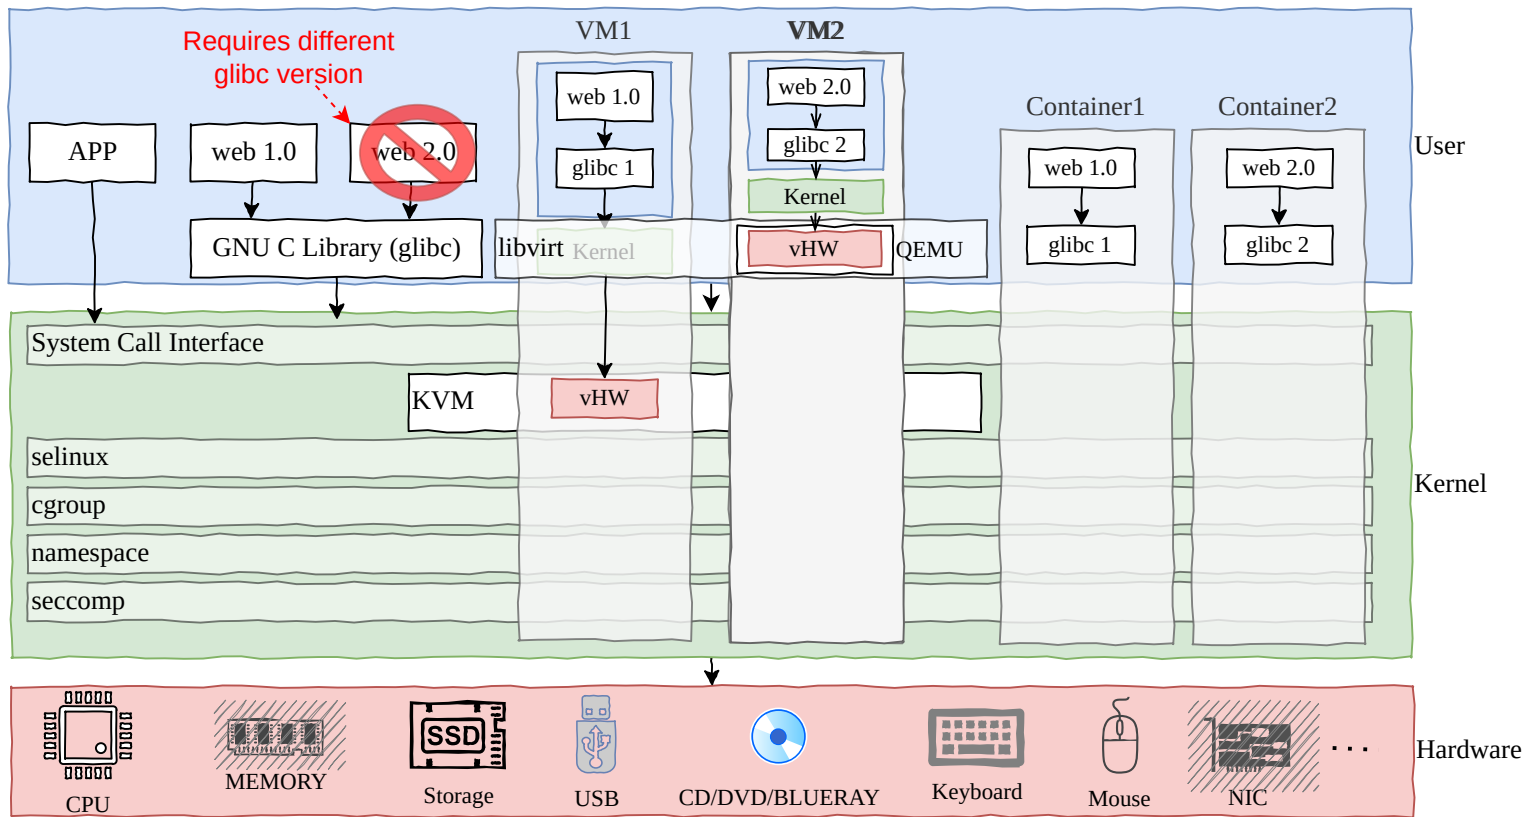
- Public hosted cluster.
- Shared resources by multiple customers.
- RH manages cluster life cycle.



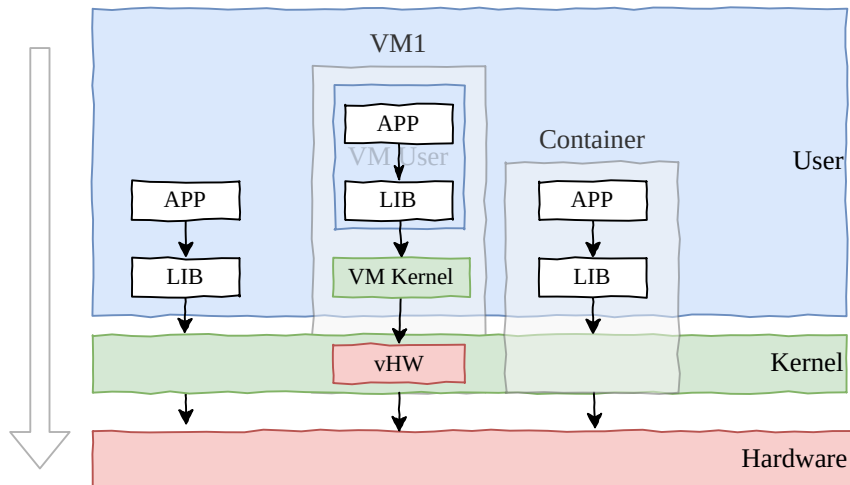
# Kubernetes Declarative Architecture



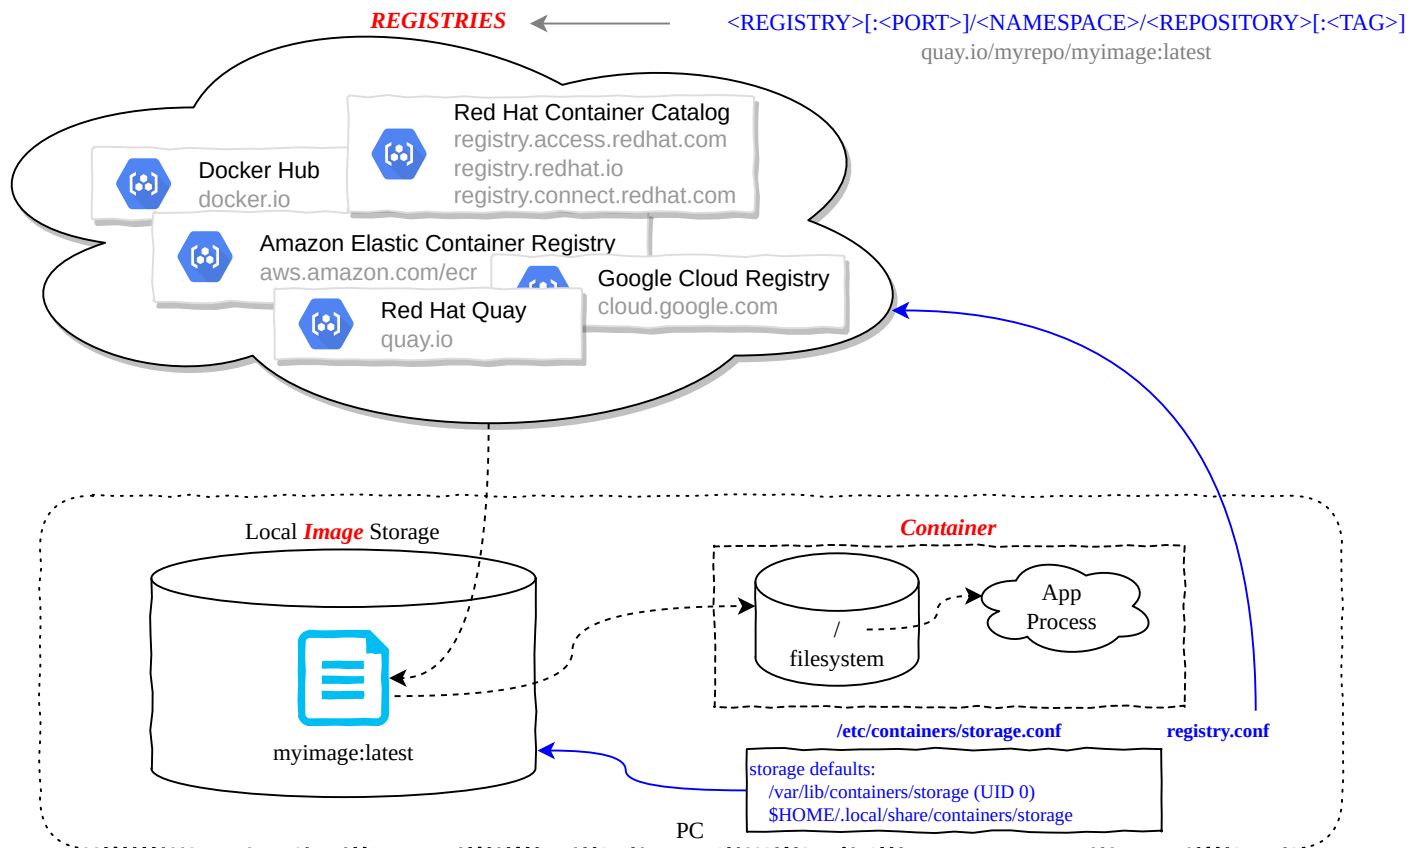
## VM vs Container



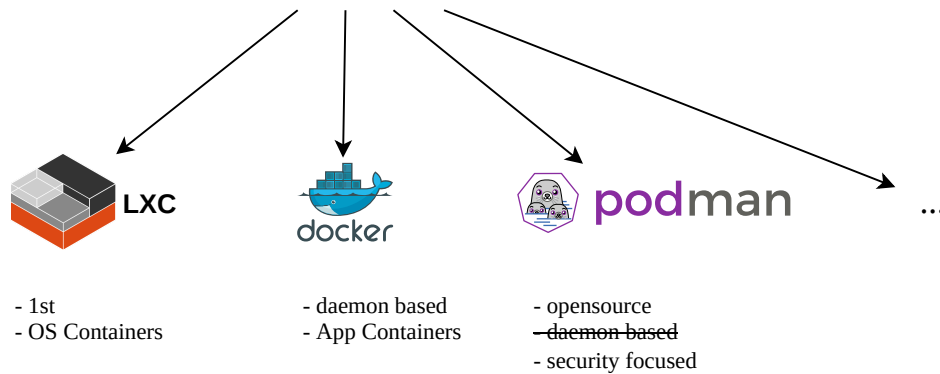
Ref: <https://www.redhat.com/en/blog/all-you-need-know-about-kvm-userspace>  
<https://www.packetcoders.io/what-is-the-difference-between-qemu-and-kvm/>




# Container Architecture



## Container Utilities



## OS Container Vs Application Containers



<https://developer.ibm.com/tutorials/multi-architecture-cri-o-container-images-for-red-hat-openshift/>

<http://www.haifux.org/lectures/299/netLec7.pdf>  
[https://kernelnewbies.org/Linux\\_2\\_X\\_XX](https://kernelnewbies.org/Linux_2_X_XX)

## Heading

man 7 namespaces

- mount (2.4.19) - 3/8/2002 -  
CAP\_SYS\_ADMIN

- pid (2.6.24) - 24/1/2008

- net (2.6.29) - 23/3/2009

- ipc (2.6.19) 29/11/2006

- uts (2.6.19) 29/11/2006

- user (3.8) 18/2/2013 no cap

- cgroup (4.6) 15/5/2016

- time - 3/2020

```
hostname  
unshare -u  
hostname abc  
hostname  
exit
```

# Podman

## Image and Registry Operations

<b>podman login</b> [-u <i>USER</i> ] [-p <i>PASS</i> ] [ <i>REGISTRY</i> ]	Only if required. Accessing private repo or updating image.
<b>podman logout</b> {-a   <i>REGISTRY</i> }	Logout of registry (-a for all).
<b>podman images</b> [-q]	List local images (-q only show id).
<b>podman rmi</b> <i>IMAGE...</i>	Removes local image(s). Use -af with caution.
<b>podman search</b> <i>KEYWORD</i>	Search registry for an image.
<b>podman pull</b> <i>SOURCE</i>	Pull image from a registry.

Where,

*SOURCE*                      [*REGISTRY*[:*PORT*]/*NAMESPACE*/]*IMAGE*[:*TAG*]  
                                  *dir:PATH*  
                                  *docker-archive:PATH*  
                                  *oci-archive:PATH*

<b>podman tag</b> <i>IMAGE</i> [: <i>TAG</i> ] <i>TARGET_NAME</i> [: <i>TAG</i> ]	Add an additional name to a local image
<b>podman push</b> <i>IMAGE</i>	Upload an image to the registry

## Container Operations

**podman run** [--name *NAME*] [-p *PORT\_INFO*] [-v *VOL\_INFO*] [-d] [-it] *IMAGE* [*CMD\_INFO*]

Where,

--name <i>NAME</i>	Container name. Autogenerated if not provided.
-p <i>PORT_INFO</i>	[ <i>LOCAL_IP</i> : ] <i>LOCAL_PORT</i> [ [ : <i>CONT_IP</i> ] : <i>CONT_PORT</i> ] Mapping between local IP:PORT to container IP:PORT
-v <i>VOL_INFO</i>	<i>LOCAL_DIR</i> : <i>CONT_DIR</i> Mapping between local dir to container dir.
-d	Run in detached mode (background).
-it	-i keep stdin open, -t allocate a pseudo-tty.
<i>IMAGE</i>	Image used to create the container.
<i>CMD_INFO</i>	<i>CMD</i> [ <i>ARG...</i> ] Command to run in container.

<b>podman exec</b> [-it] <i>CONTAINER</i> <i>CMD_INFO</i>	Execute command inside running container.
<b>podman ps</b> [-a] [-q]	List containers (-a for all, -q only show container id).
<b>podman rm</b> <i>CONTAINER...</i>	Remove one or more stopped containers. (-f includes running and paused containers).
<b>podman start stop restart</b> <i>CONTAINER...</i>	Start, stop or restart one or more containers.
<b>podman kill</b> [-s <i>SIGNAL</i> ] <i>CONTAINER...</i>	Send signal to one or more containers.

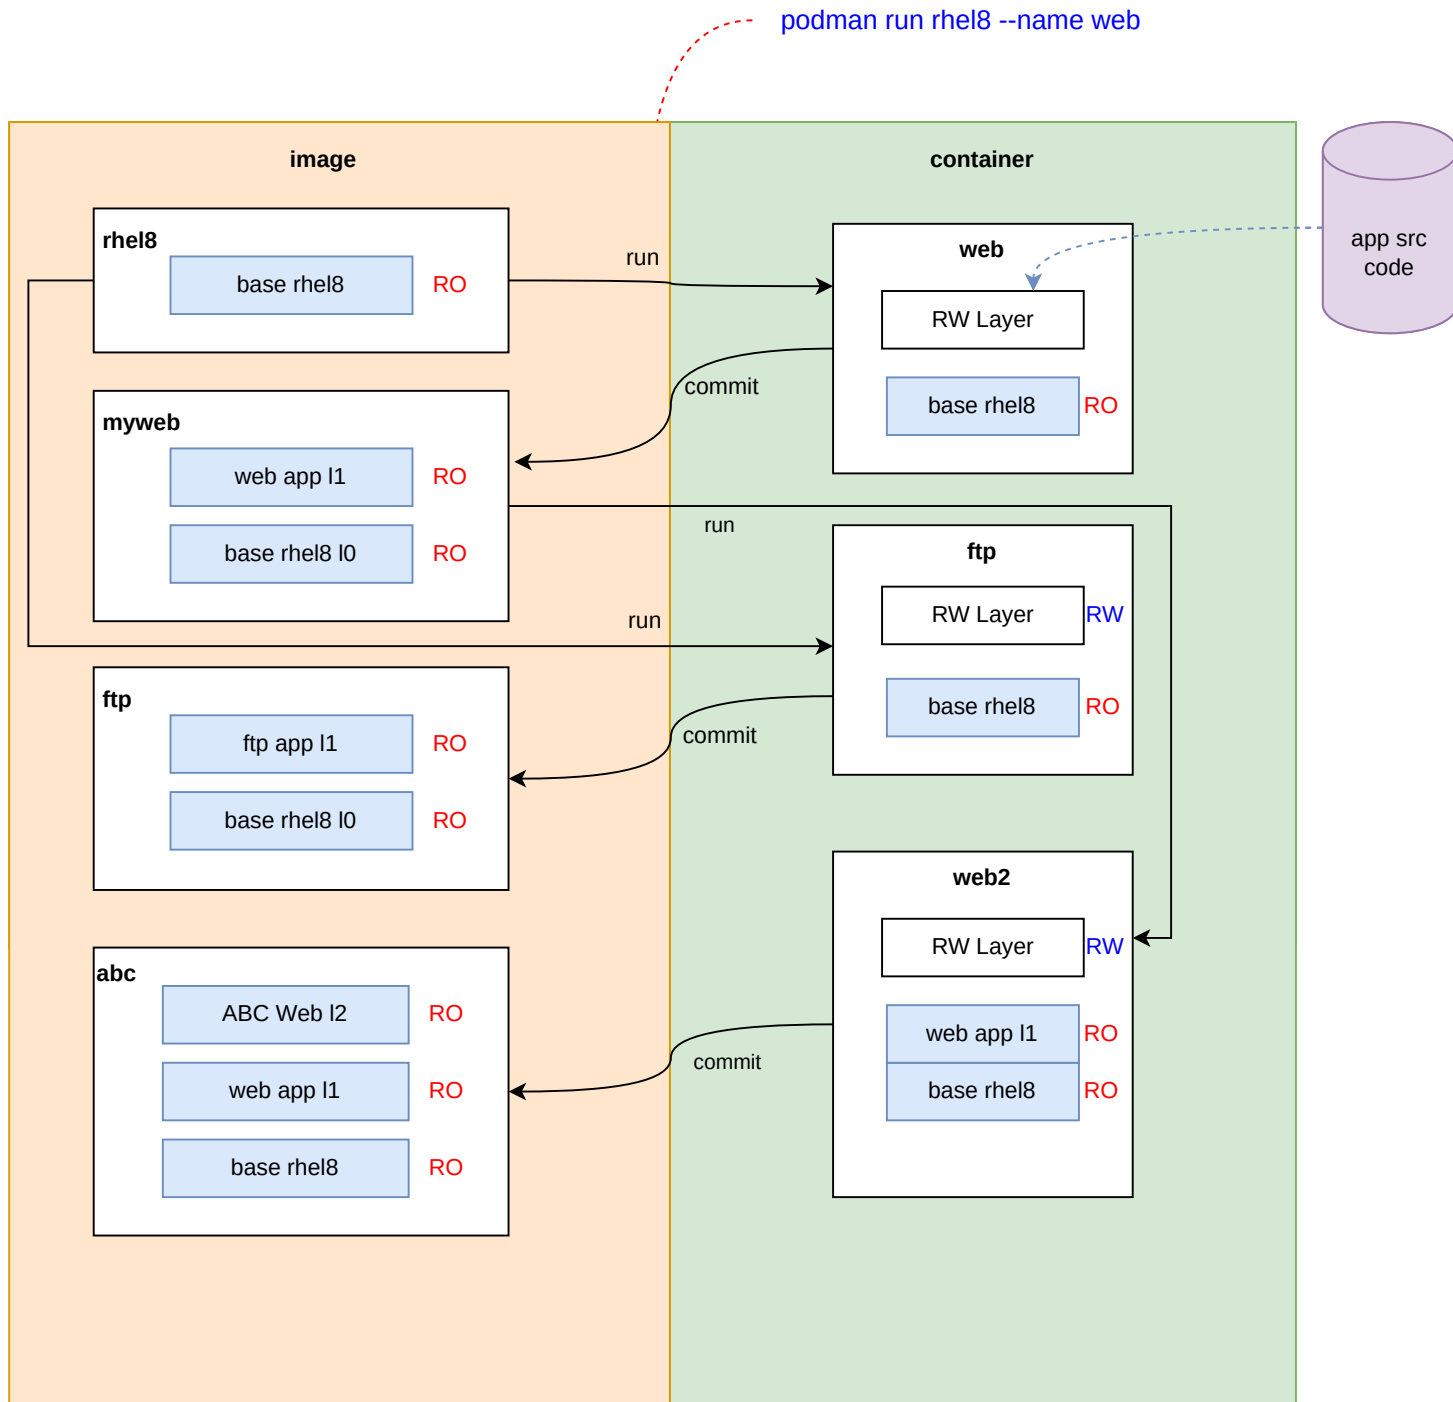
For more info:

**podman --help** OR **man podman**

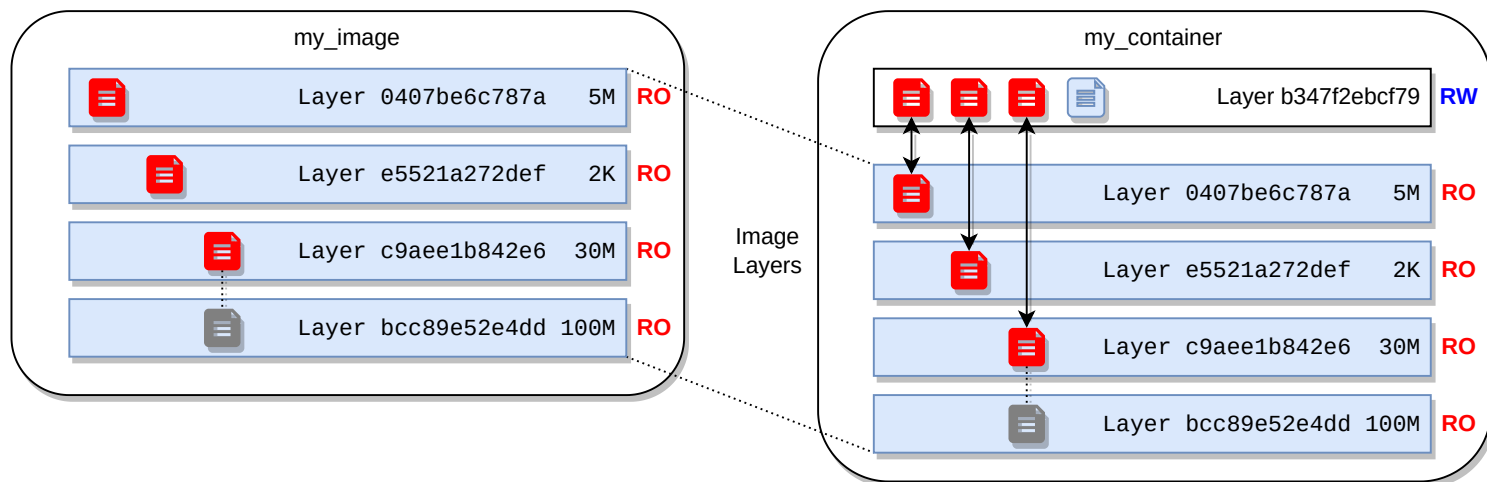
Each sub command has it's own man page. i.e man podman-run, man podman-images, etc.

# Creating Image

1. Manual
2. Dockerfile/Containerfile
3. Source-To-Image(s2i/STI)
  - a) get runtime image and create container
  - b) clone source code into container
  - c) compile source code
  - d) deploy/publish compiled app
  - e) cleanup
  - f) save container as image

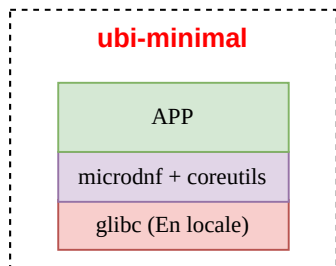


# UnionFS - A Stackable Unification File System



## BASE IMAGE TYPES


### MINIMAL



**Designed for apps that contain their own dependencies (Python, Node.js, .NET, etc.)**

- Minimized pre-installed content set
- no suid binaries
- minimal pkg mgr (install, update & remove)


### PLATFORM



**For any apps that runs on RHEL**

- Unified, OpenSSL crypto stack
- Full YUM stack
- Includes useful basic OS tools (tar, gzip, vi, etc)

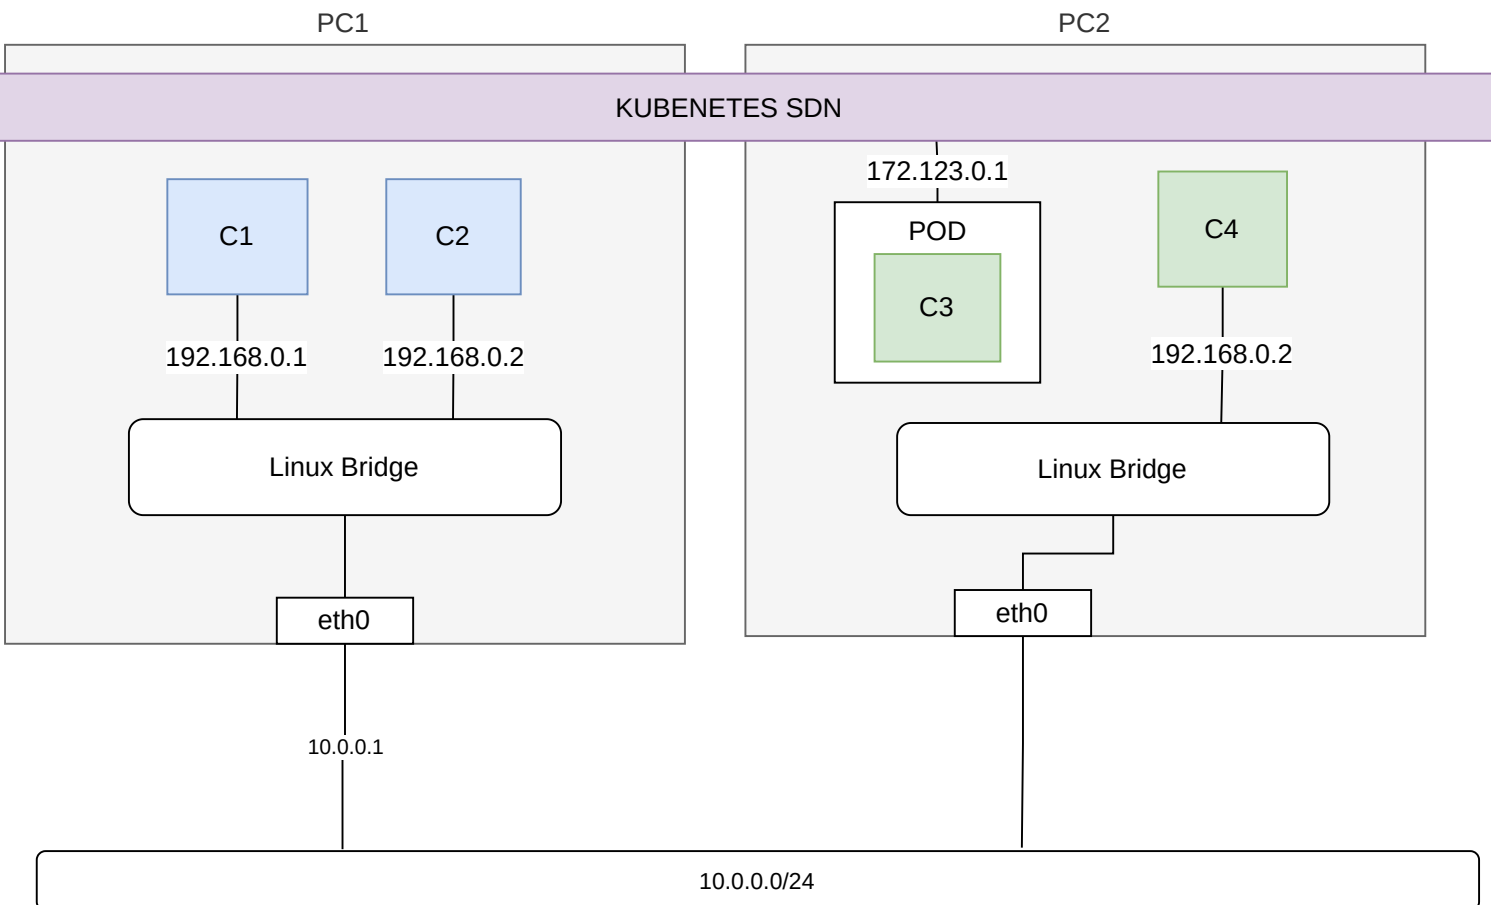
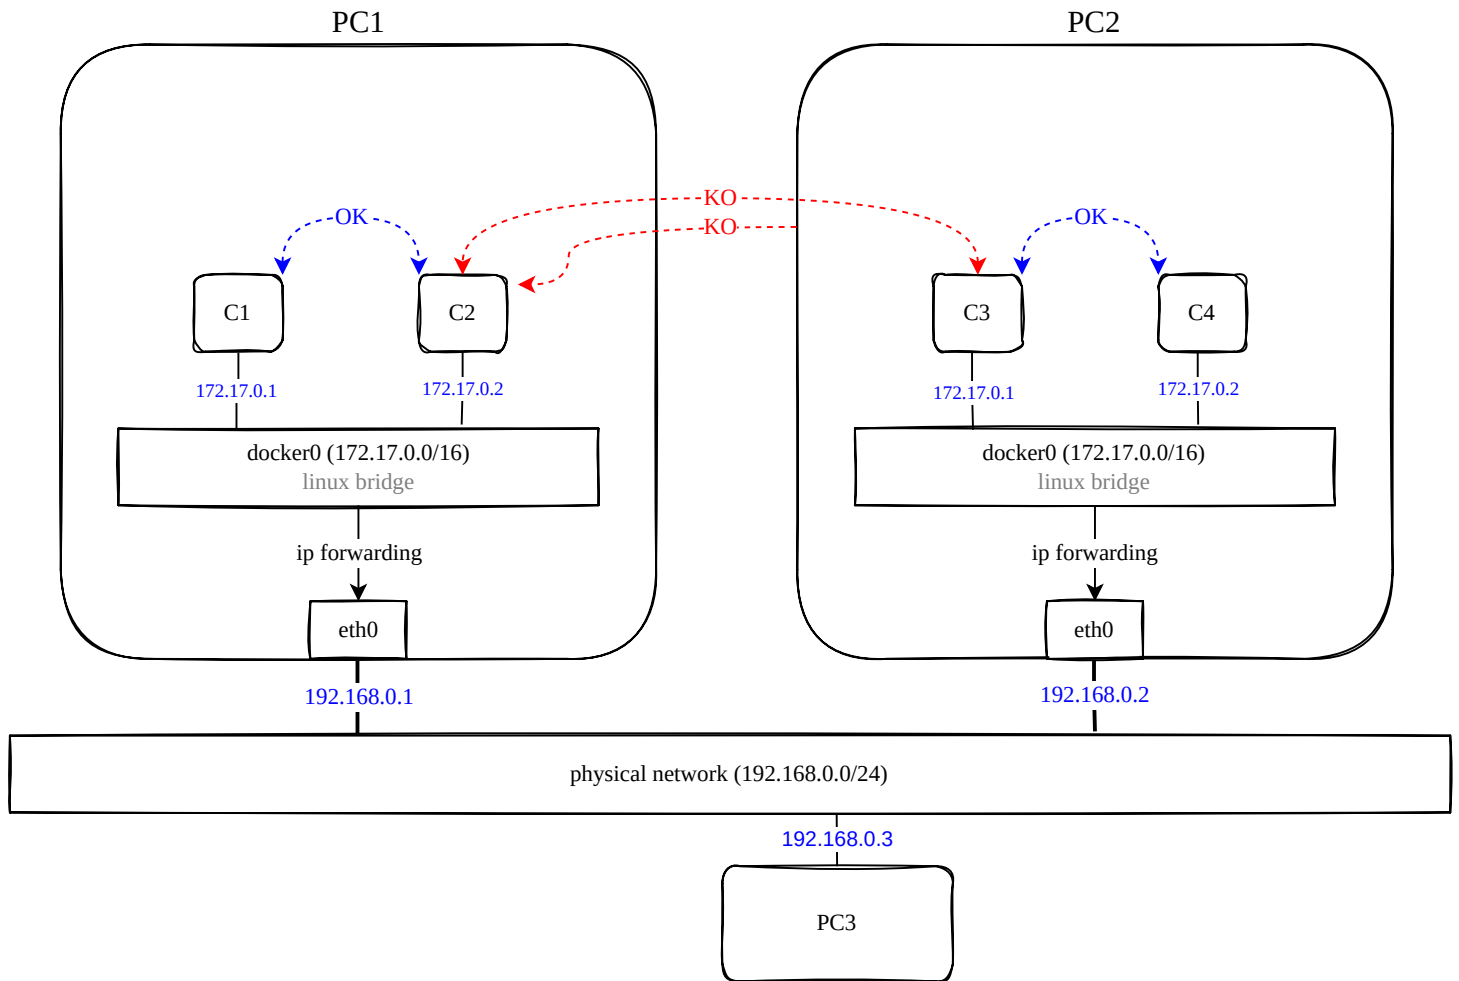
### MULTI-SERVICE



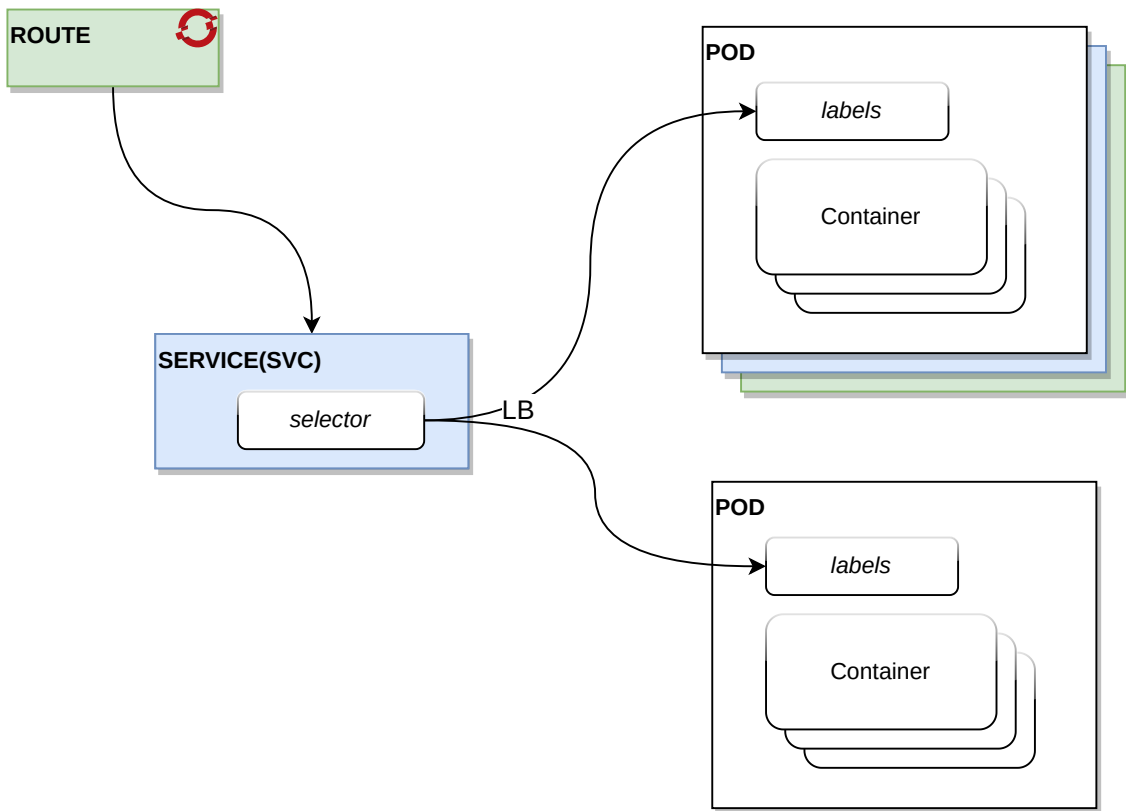
**Eases running multi-service in single container**

- configured to run systemd on start
- allows you to enable th services at build time





## Route, Service and Pod Relationship



### POD

A pod contains one or more containers.

### SERVICE

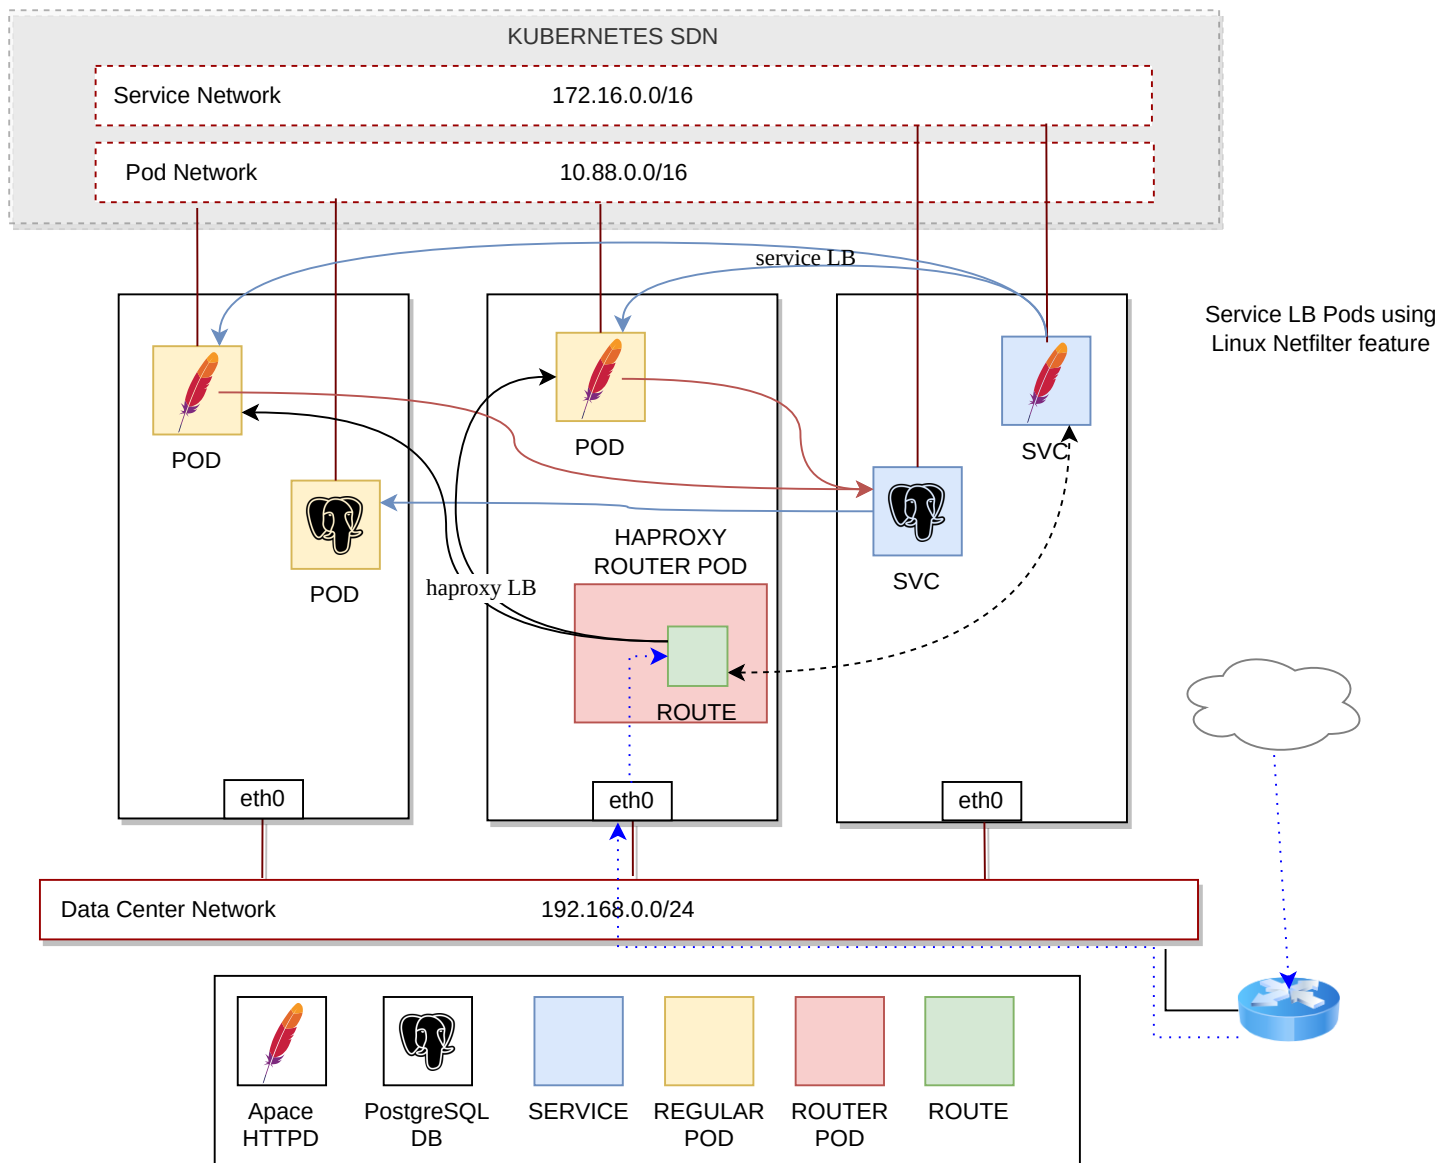
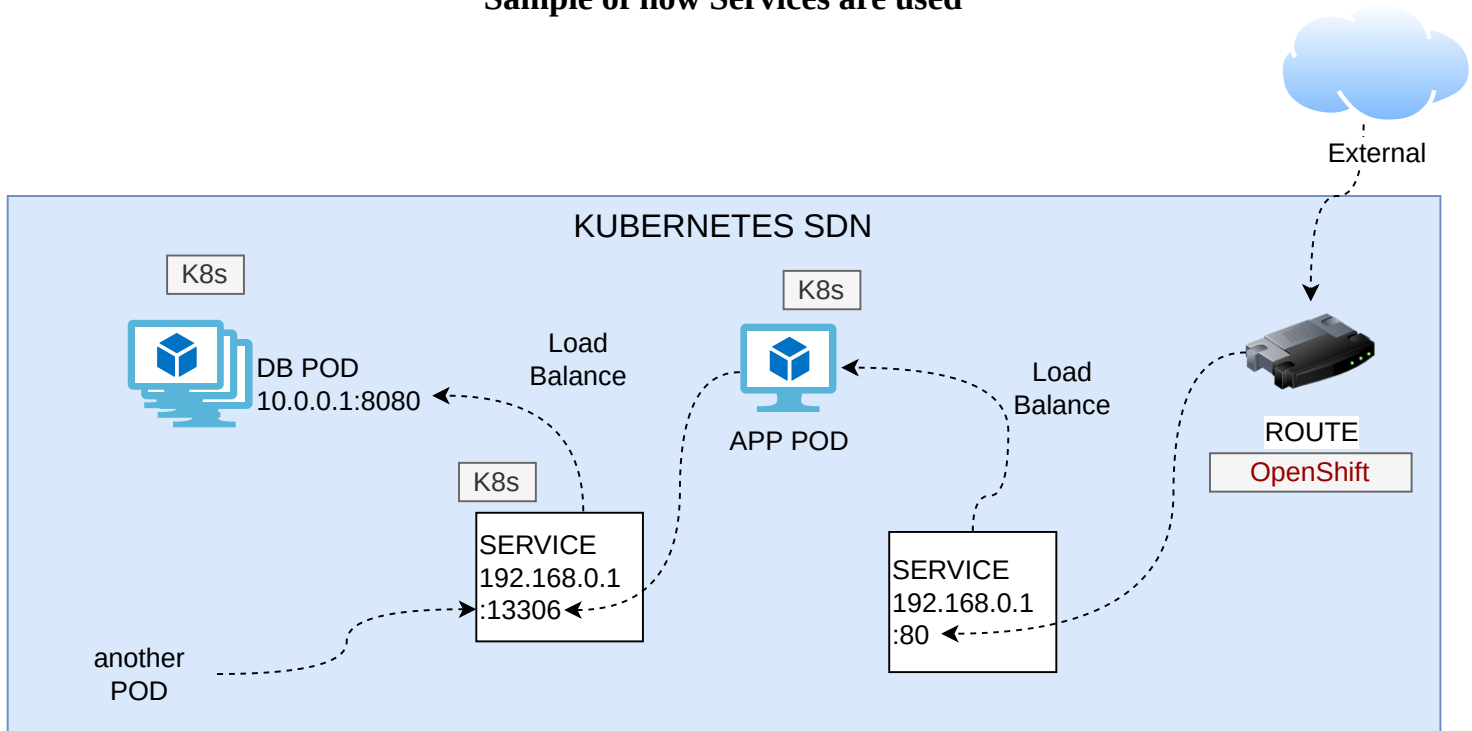
A service references the pod(s) by using the label selector.  
The service load balances the connections between all the pods.

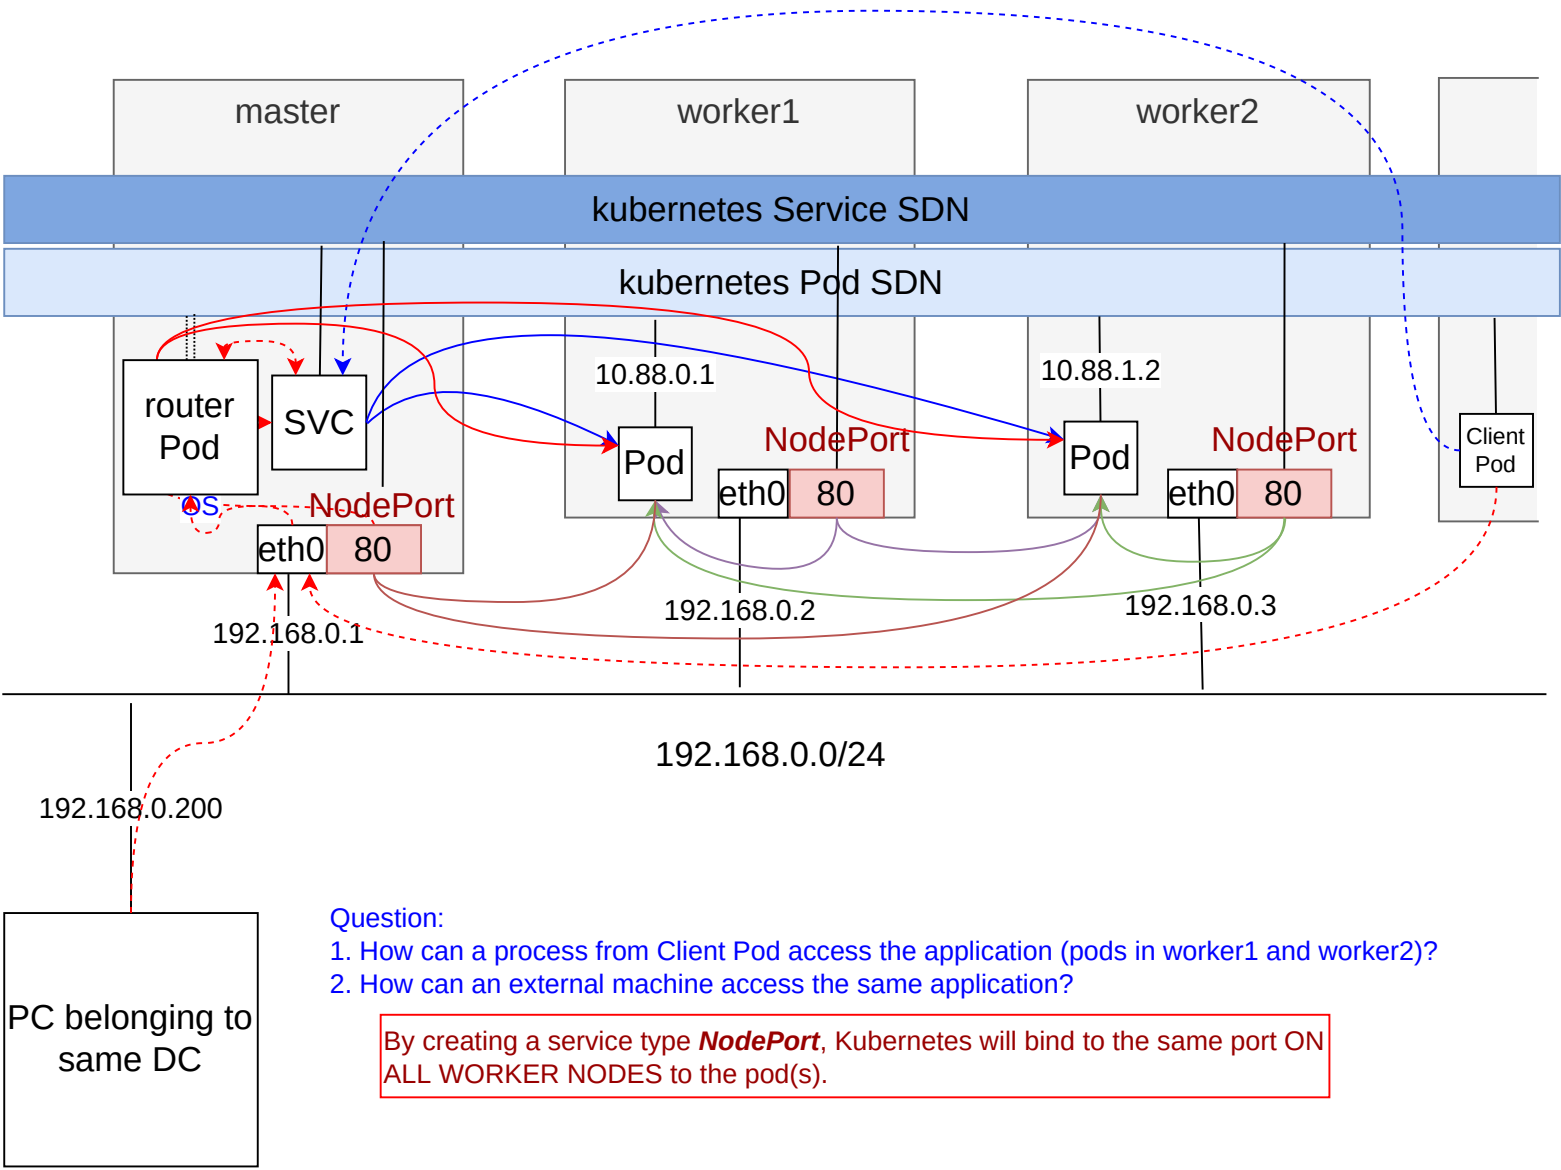
### ROUTE

A route exposes the service to the external world.

**Warning:** A service "can" refer to different pods, if the pods have the same label.

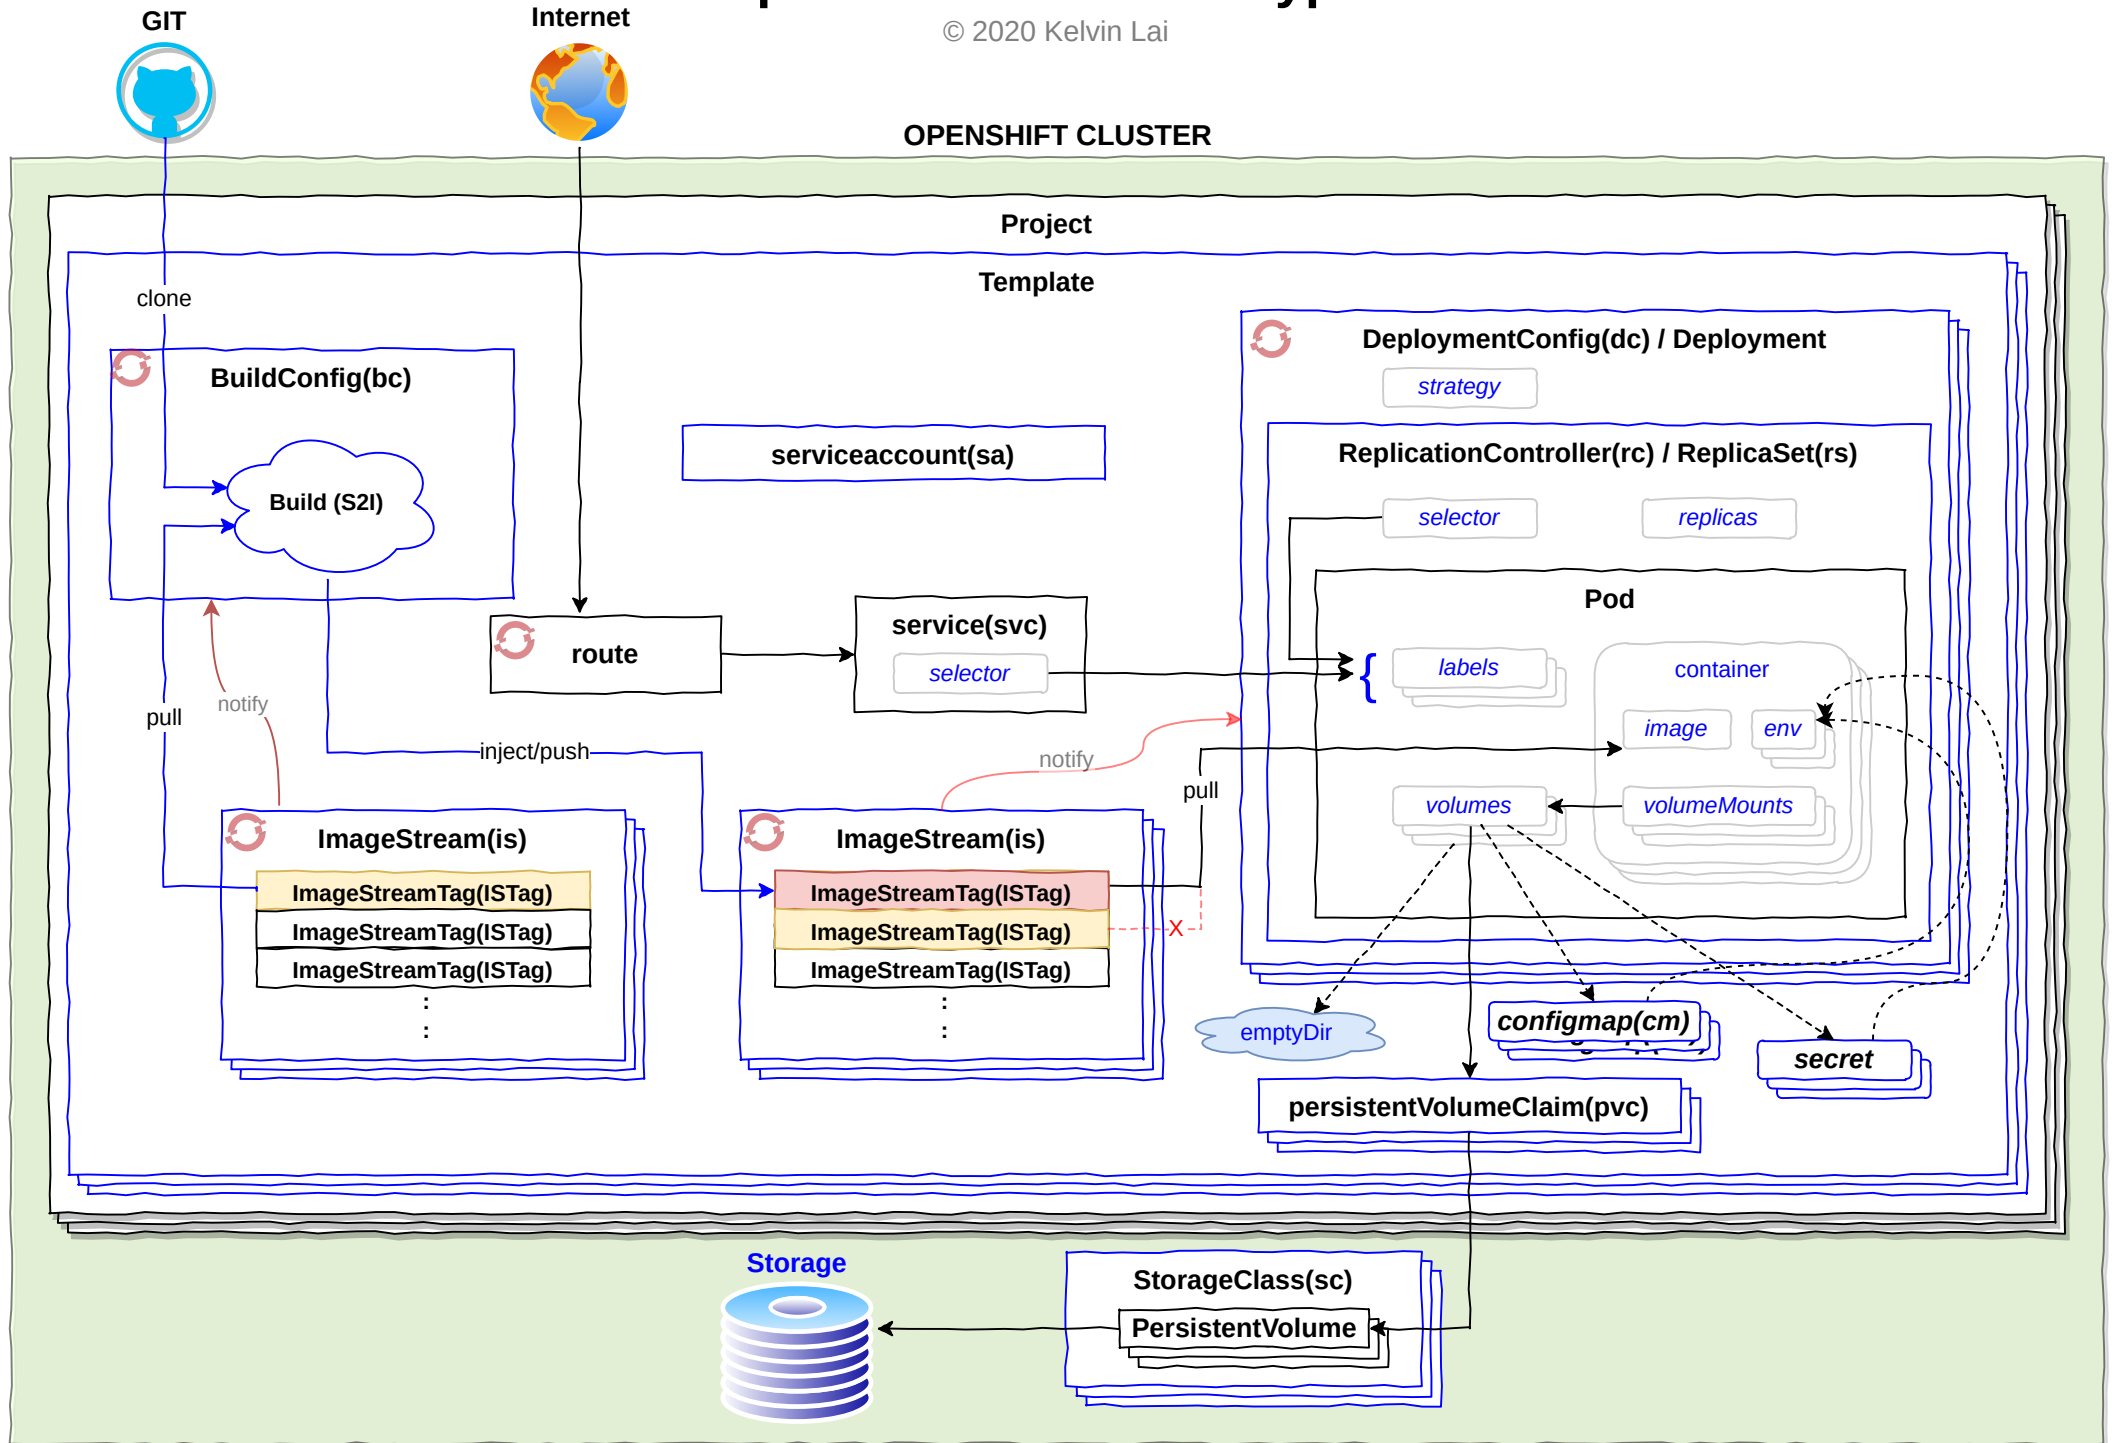
## Sample of how Services are used





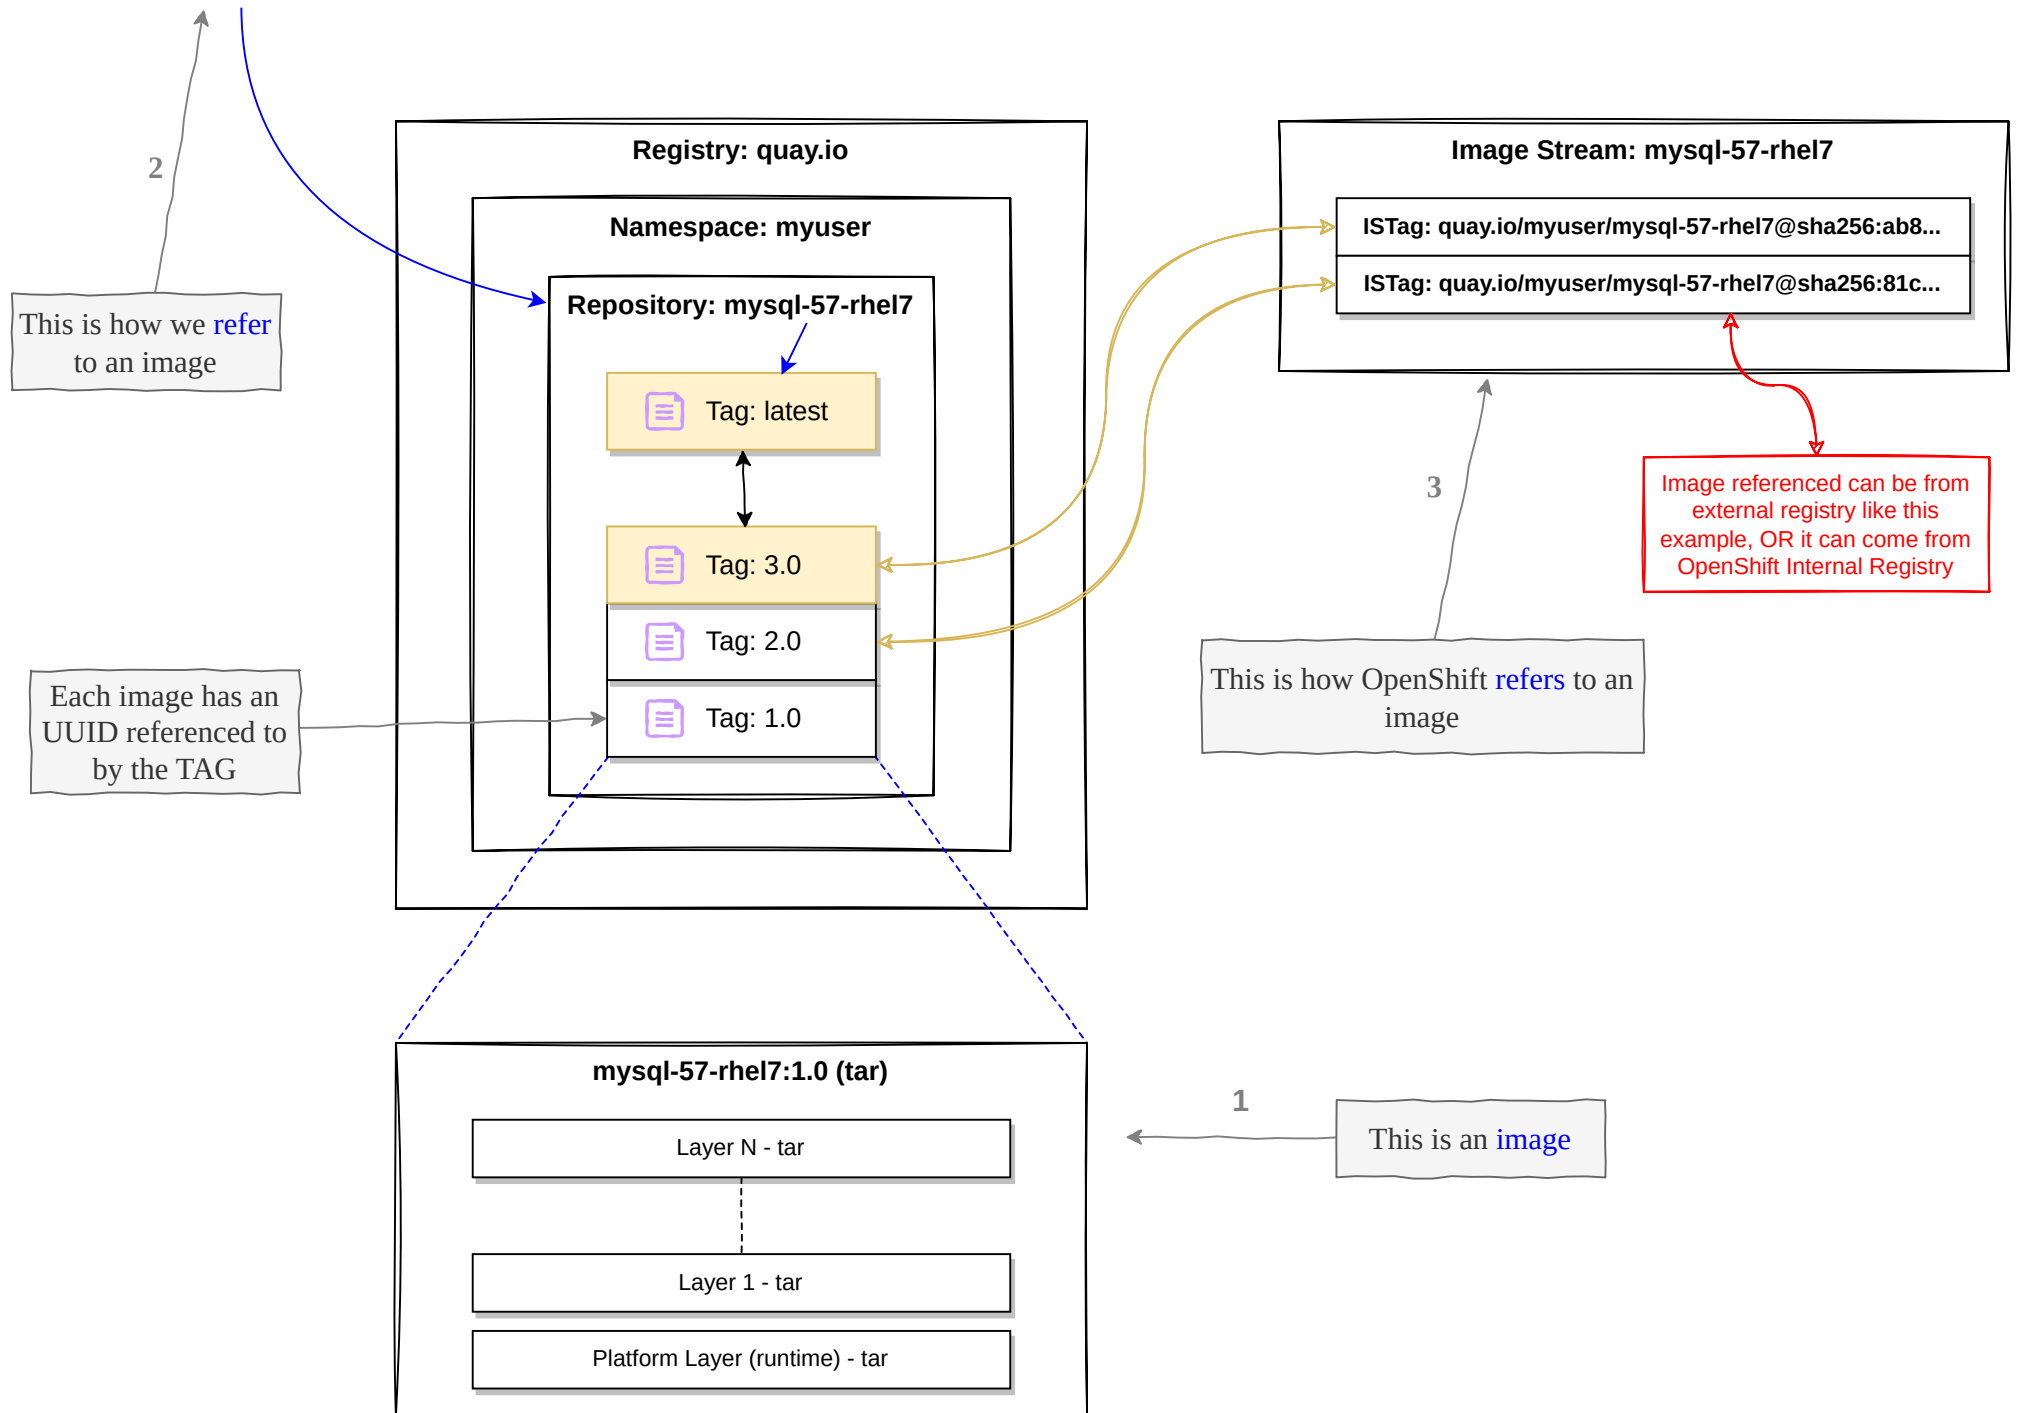
# OpenShift Resource Types

© 2020 Kelvin Lai



Container Image Naming Convention: <REGISTRY>[:<PORT>]/<NAMESPACE>/<REPOSITORY>[:<TAG>]

podman pull [quay.io/myuser/mysql-57-rhel7](#)



# Deploying Applications with OpenShift

Methods to create applications:

1. Using existing containerised applications

```
oc new-app --docker-image=<IMAGE>
```

2. From Source Code using S2I

```
oc new-app <URL>
```

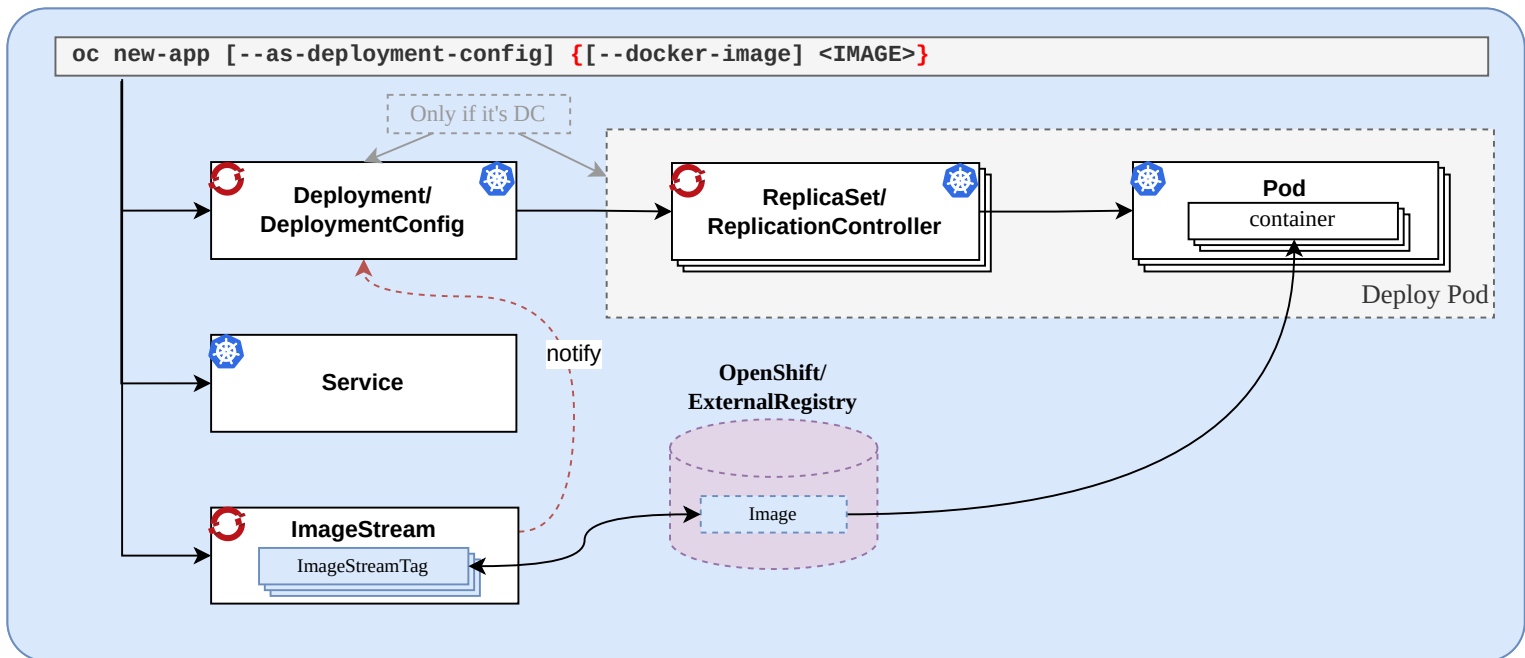
3. Using yaml/json file

```
oc new-app -f <FILE>.yaml
```

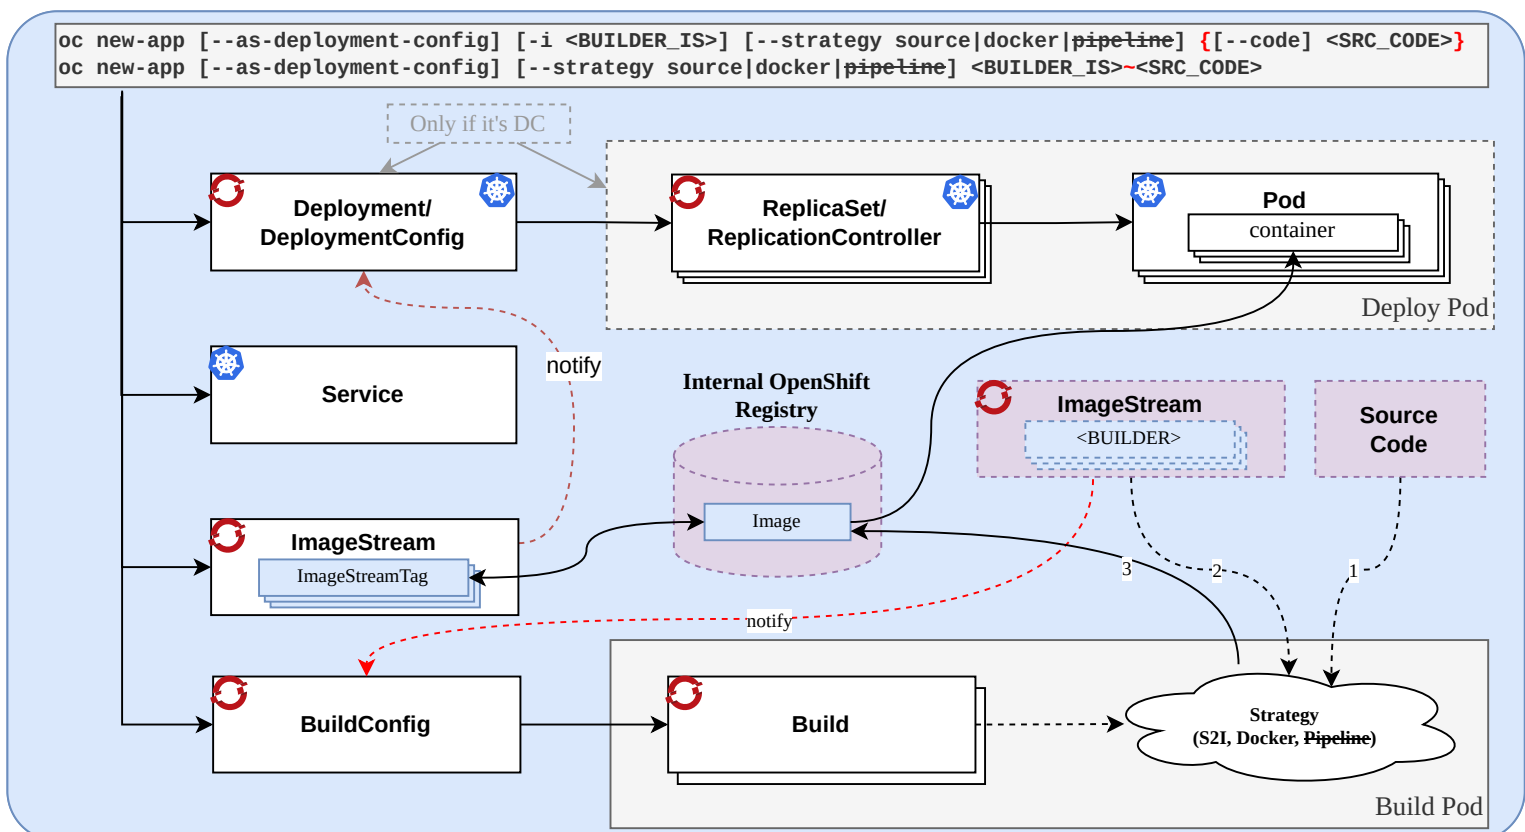
4. Using template

```
oc new-app --template=<TEMPLATE> --param=<PARAM> --param-file=<PARAM_FILE>
```

## 1. Use Existing Image



## 2. Managed Life Cycle



```
oc new-app -i myphp https://github.com/user/myapp#branch --context-dir <DIR>
```

```
oc new-app -i myphp:7.1 https://github.com/user/myapp
```

```
oc new-app myphp:7.1~https://github.com/user/myapp
```

NOTE: -i option needs git client to be installed

## Options

**-o json|yaml** inspect resource definitions without creating

**--name <NAME>** adds a label "app=<NAME>" to all resources, Use `oc delete all -l "app=<NAME>"` to cleanup

## IMPORT IMAGES

```
oc import-image <IMG_STREAM> [--confirm] --from <IMAGE> [--insecure]
where,
    <IMAGE> = <REGISTRY>[:<PORT>]/<NAMESPACE>/<REPOSITORY>[:<TAG>]
```

oc new-app command in OpenShift 4.5 makes use of [deployment](#) resource. Use --as-deployment-config if you wish to create [deployment config](#) instead.

## SERVICE(SVC)

```
oc expose <DC/DEPLOYMENT/RC/RS/POD> <RESOURCE_NAME>
```

```
DNS NAME = <SVC>.<PROJ>[.svc.cluster.local]
ENVIRONMENT VARIABLE IN POD = <SVC>_SERVICE_HOST
```



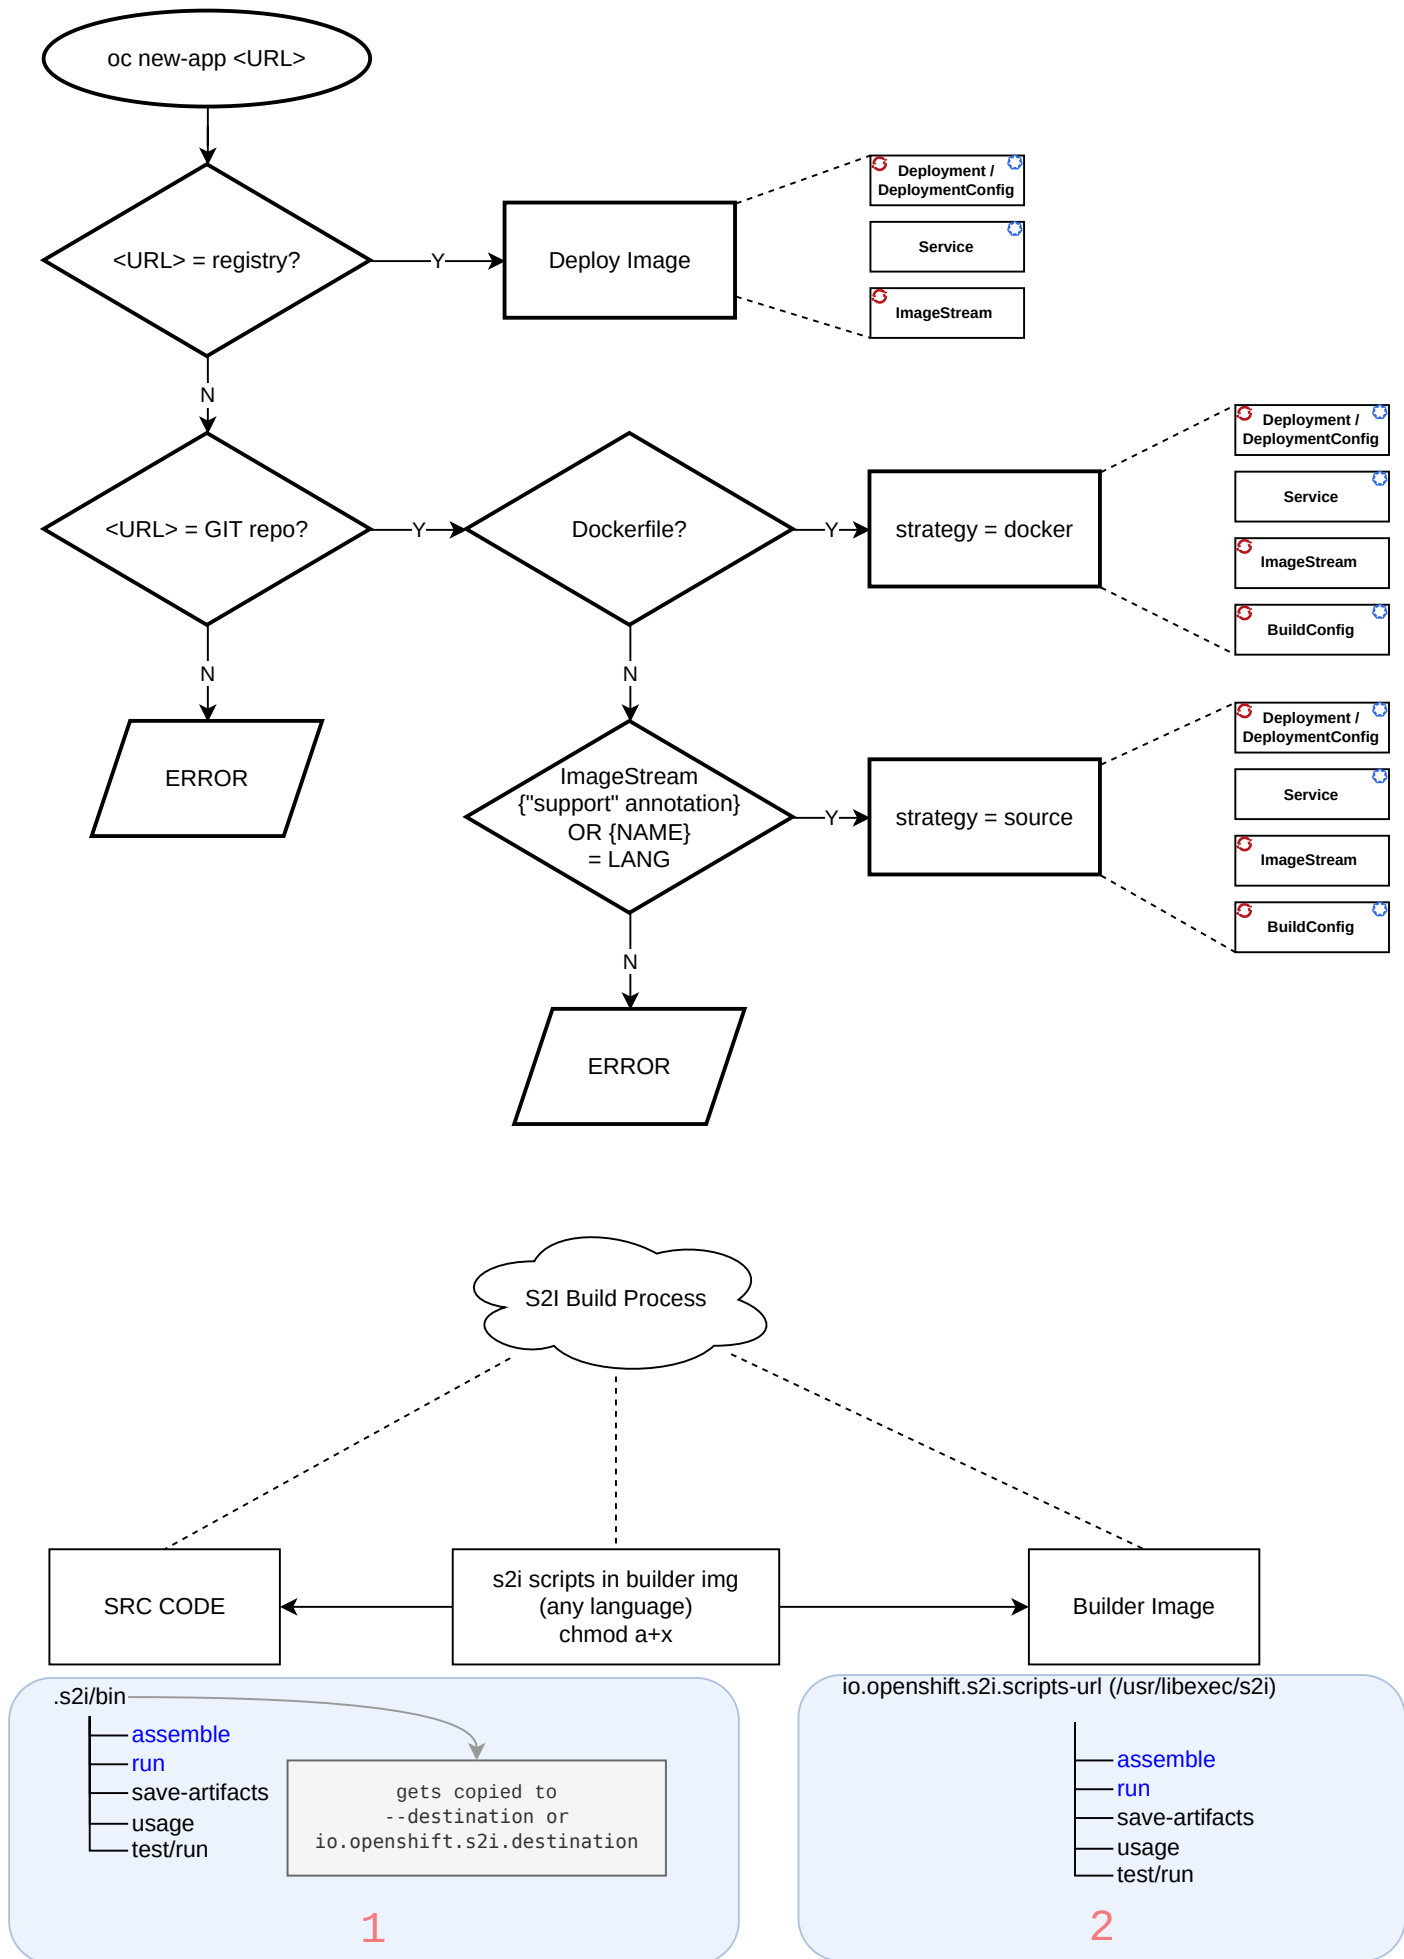
## ROUTE

```
oc expose svc <SVC_NAME> [--name <ROUTE_NAME>] [--hostname <FQDN>]
```

```
DNS DEFAULT NAME = <ROUTE_NAME>.<PROJ>.<DOMAIN WILDCARD>
<DOMAIN_WILDCARD> = apps.<BASE_DOMAIN>
```



# oc new-app flowchart



# Template

**apiVersion:** template.openshift.io/v1

**kind:** Template

**metadata:**

name: mytemplate

annotations:

description: "Description"

**objects:**

- apiVersion: v1

kind: Pod

metadata:

name: \${APP\_NAME}

spec:

containers:

- env:

- name: ACCESS\_CODE

value: \${APP\_PASS}

image: superapp/hyperimage

name: myApp

ports:

- containerPort: 8080

protocol: TCP

**parameters:**

- description: Name of Pod

name: POD\_NAME

value: myPod

required: true

- description: Application Secret Access Code

name: APP\_PASS

generate: expression

from: "[a-zA-Z0-9]{8}"

**labels:**

mylabel: myapp

Object Creation Order  
↓

```

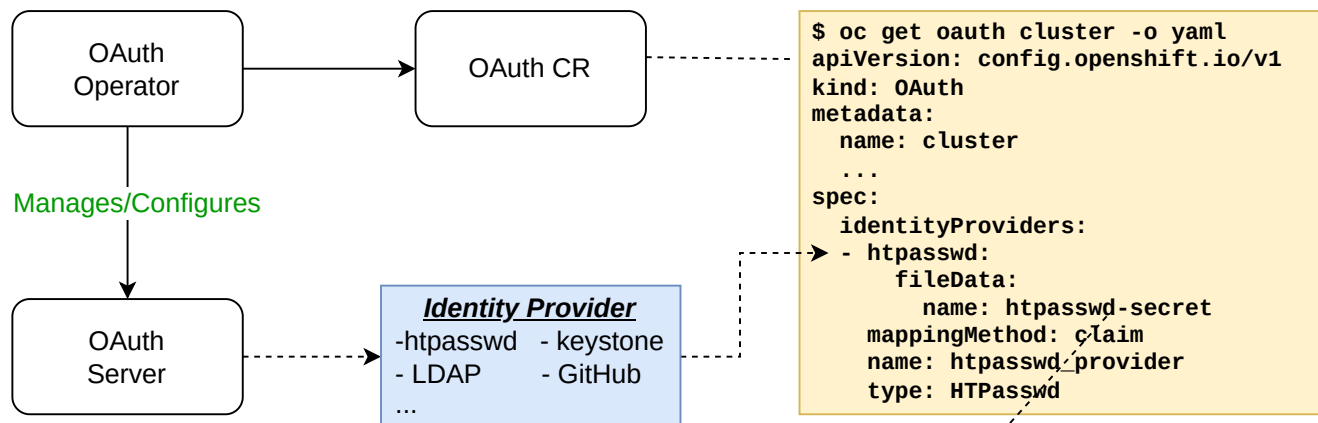

pipeline {
    options {                                // Global Options - can be overridden per stage basis
        timeout(time: 30, unit: 'MINUTES')
    }
    agent {                                  // Execution Context (where to run)
        node {
            label 'master'                  // IS used to run pipeline, (master - default minimal linux runtime)
        }
    }
    environment {                            // Global Vars
        DEV_PROJ = "super-app-v1"
        APP_NAME = "myapp"
    }
    stages {
        stage('stage 1') {
            steps {
                script {                    // DSL code (Groovy language code) must be embedded within script element
                    openshift.withCluster() {
                        openshift.withProject(env.DEV_PROJ) {
                            openshift.selector("bc", "${APP_NAME}").startBuild("--wait=true", "--follow=true")
                        }
                    }
                }
            }
        }
        stage('stage 2') {                // NOTE: project switches to default proj between stages.
            steps {
                sh 'echo hello from stage 2!'
            }
        }
        stage('manual approval') {
            steps {
                timeout(time: 60, unit: 'MINUTES') {
                    input message: "Move to stage 3?"
                }
            }
        }
        stage('stage 3') {
            steps {
                sh 'echo hello from stage 3!. This is the last stage...'
            }
        }
    }
}

```

```
$ oc get clusteroperator authentication
```

NAME	VERSION	AVAILABLE	PROGRESSING	DEGRADED	SINCE
authentication	4.5.4	True	False	False	40d

```
$ oc get crd authentications.operator.openshift.io
```



Manages/Configures

Issues Access Tokens

**Identity Provider**

- htpasswd - keystone
- LDAP - GitHub
- ...

```
$ oc get oauth cluster -o yaml
apiVersion: config.openshift.io/v1
kind: OAuth
metadata:
  name: cluster
  ...
spec:
  identityProviders:
    - htpasswd:
      fileName:
        name: htpasswd-secret
      mappingMethod: claim
      name: htpasswd_provider
      type: HTPasswd
```

```
$ oc get secret htpasswd-secret -n openshift-config -o yaml
apiVersion: v1
data:
  htpasswd: YWRT...
kind: Secret
metadata:
  name: htpasswd-secret
namespace: openshift-config
resourceVersion: "90940"
type: Opaque
```

# OPENSIFT CLUSTER (DO380-Trusted TLS Certs)

1 `oc create configmap <CONFIGMAP> --from-file ca-bundle.crt=<CERTIFICATE> -n openshift-config`

Namespace: *openshift-config*

ConfigMap: **combined-certs**

data  
combined-cert.pem

Secret: **custom-tls**

data  
: combined-cert.pem  
: wildcard-api-key.pem

*cluster*

trustedCA  
combined-certs

2 `oc patch proxy/cluster --type=merge --patch='{"spec":{"trustedCA":{"name":"<CONFIGMAP_NAME>"}}}'`

3 `oc create secret tls <SECRET_NAME> --cert <CERTIFICATE> --key <KEY> -n <NAMESPACE>`

Namespace: *openshift-ingress*

Secret: **custom-tls-bundle**

data  
: combined-cert.pem  
: wildcard-api-key.pem

Namespace: *openshift-ingress-operator*

ingresscontrollers.operator: **default**

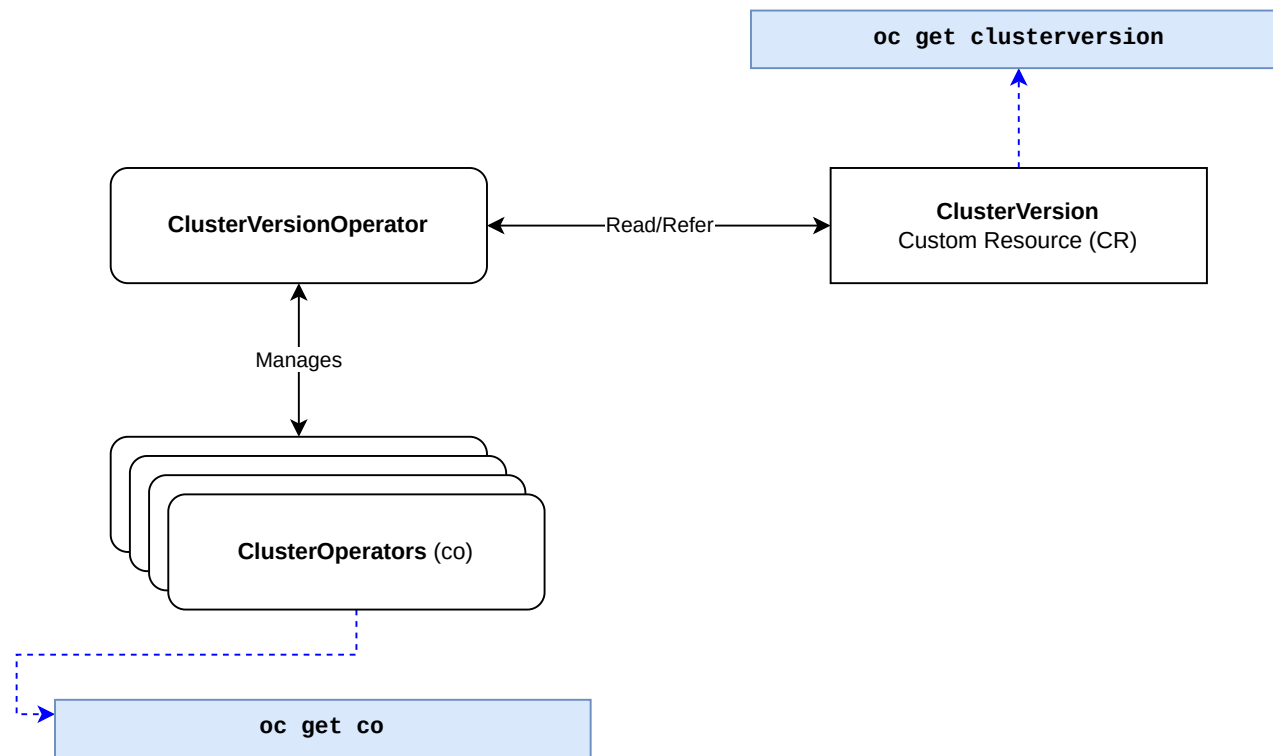
defaultCertificate  
custom-tls-bundle

*cluster*

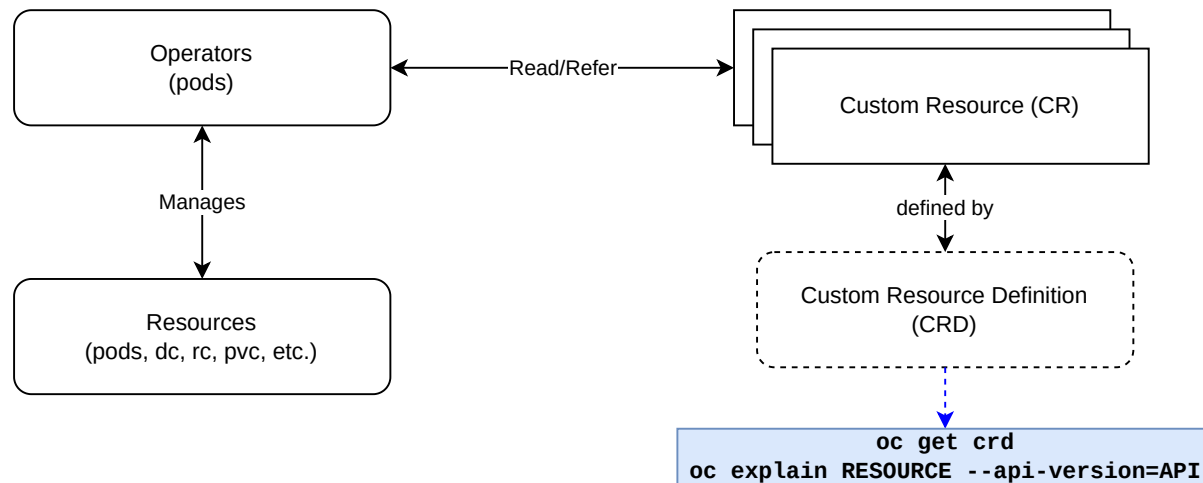
servingCerts  
api.ocp4.example.com  
custom-tls

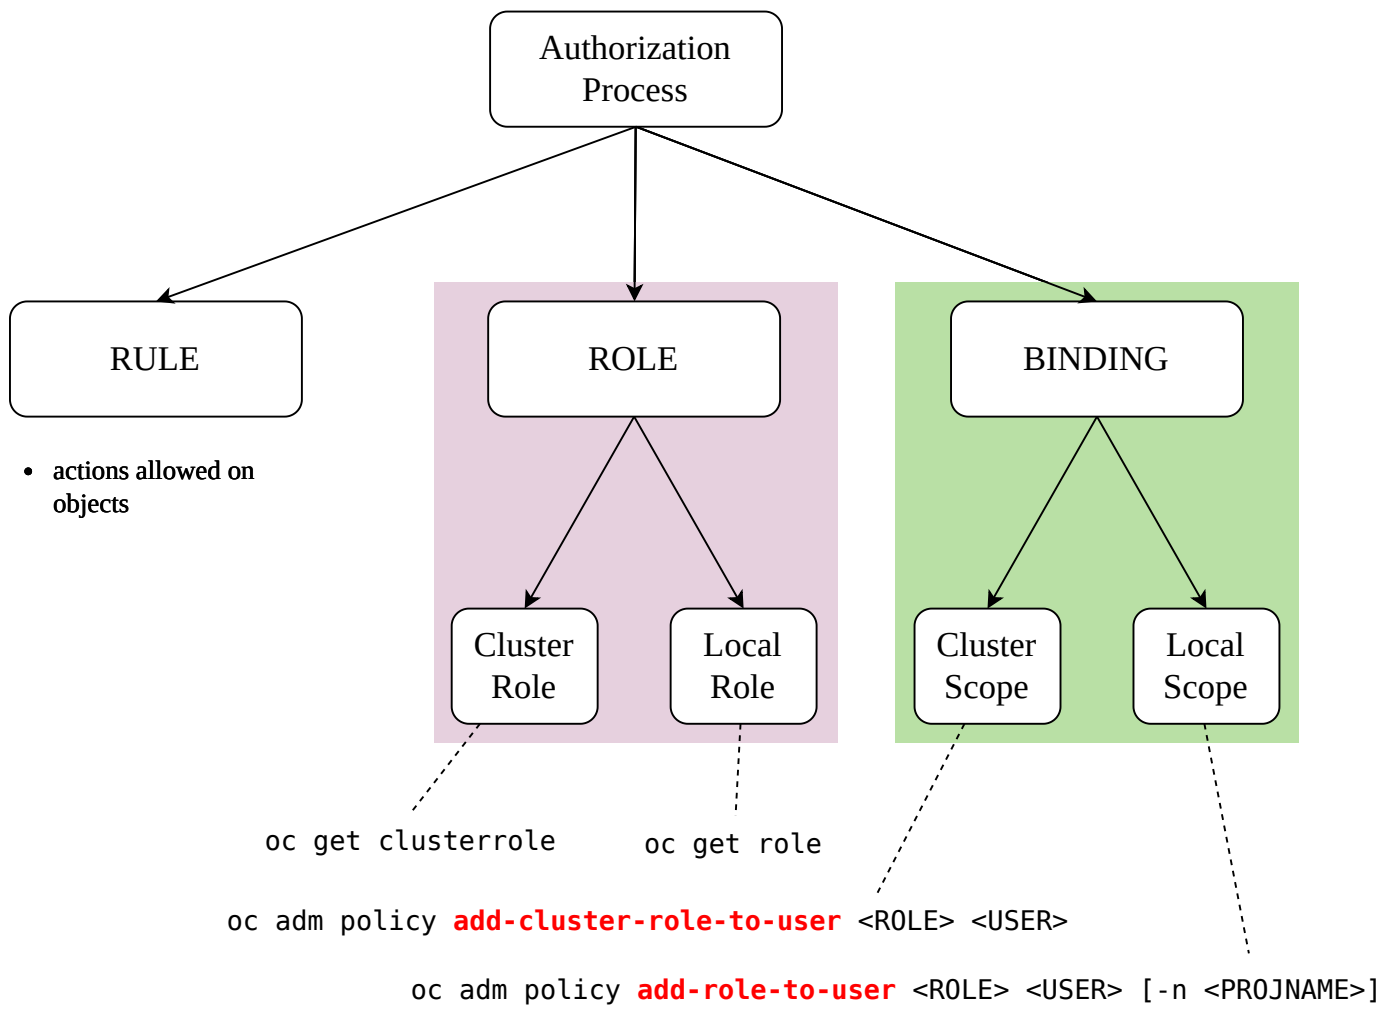
4 `oc patch ingresscontrollers.operator/default --type=merge --patch='{"spec":{"defaultCertificate":{"name":"<SECRET>"}}}' -n openshift-ingress-operator`

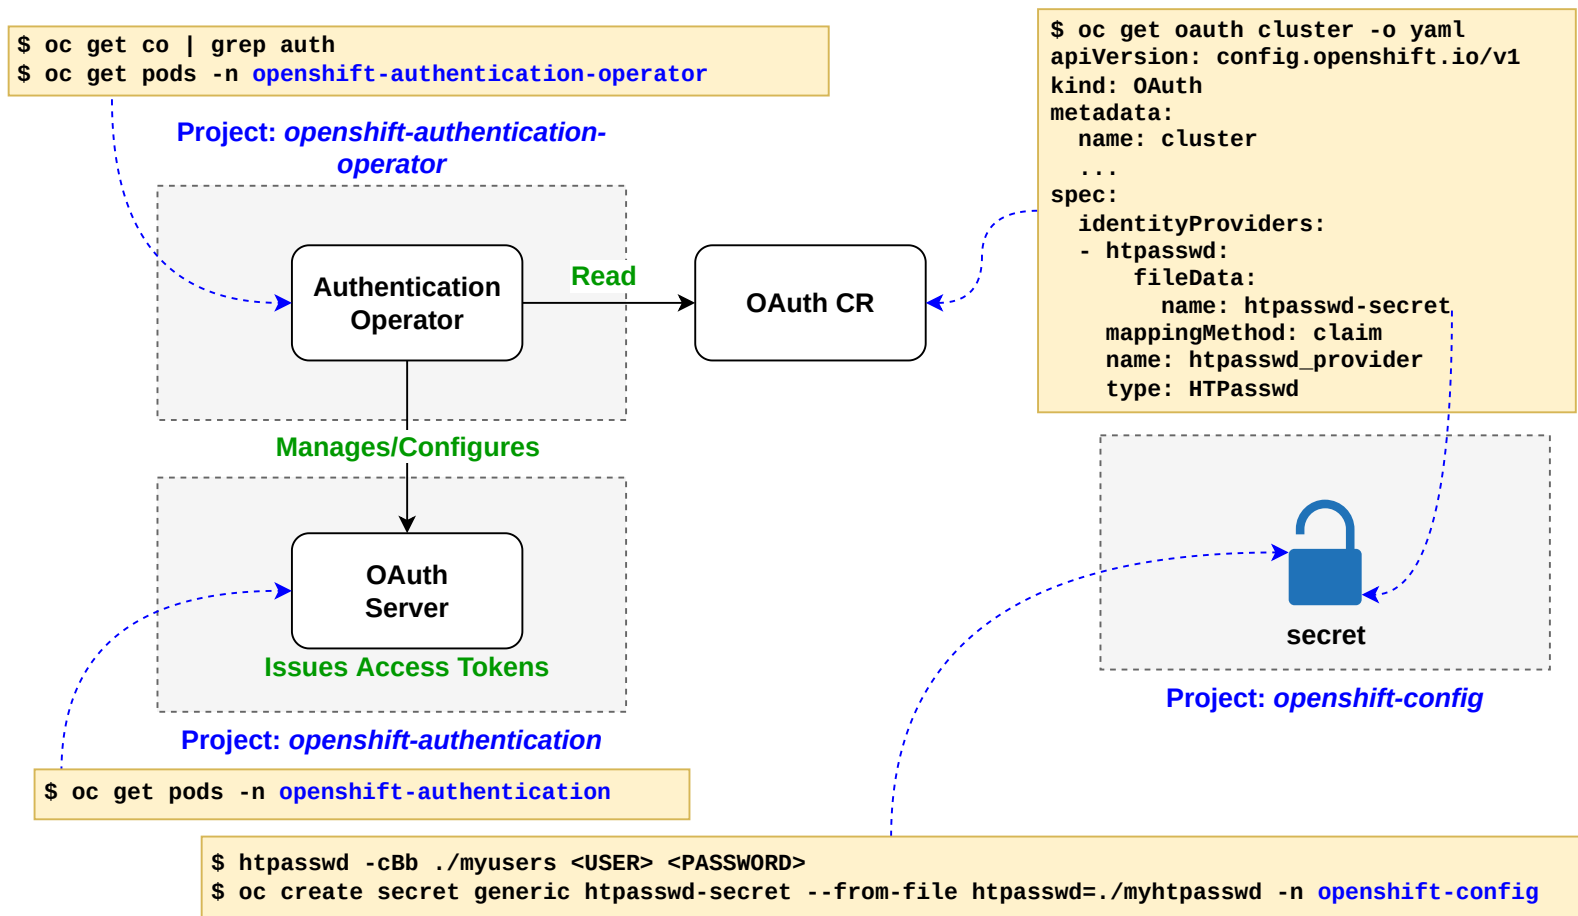
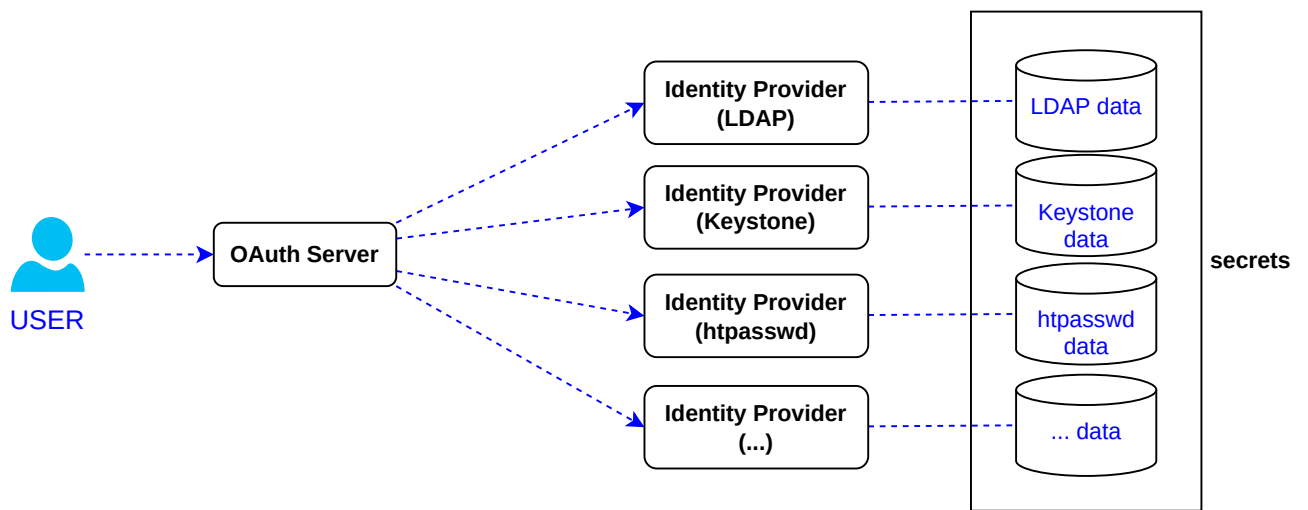
5 `oc patch apiserver cluster --type=merge --patch='{"spec":{"servingCerts":{"namedCertificates":[{"names":["<API-SERVER-URL>"],"servingCertificate":{"name":"<SECRET-NAME>"}}]}}}'`



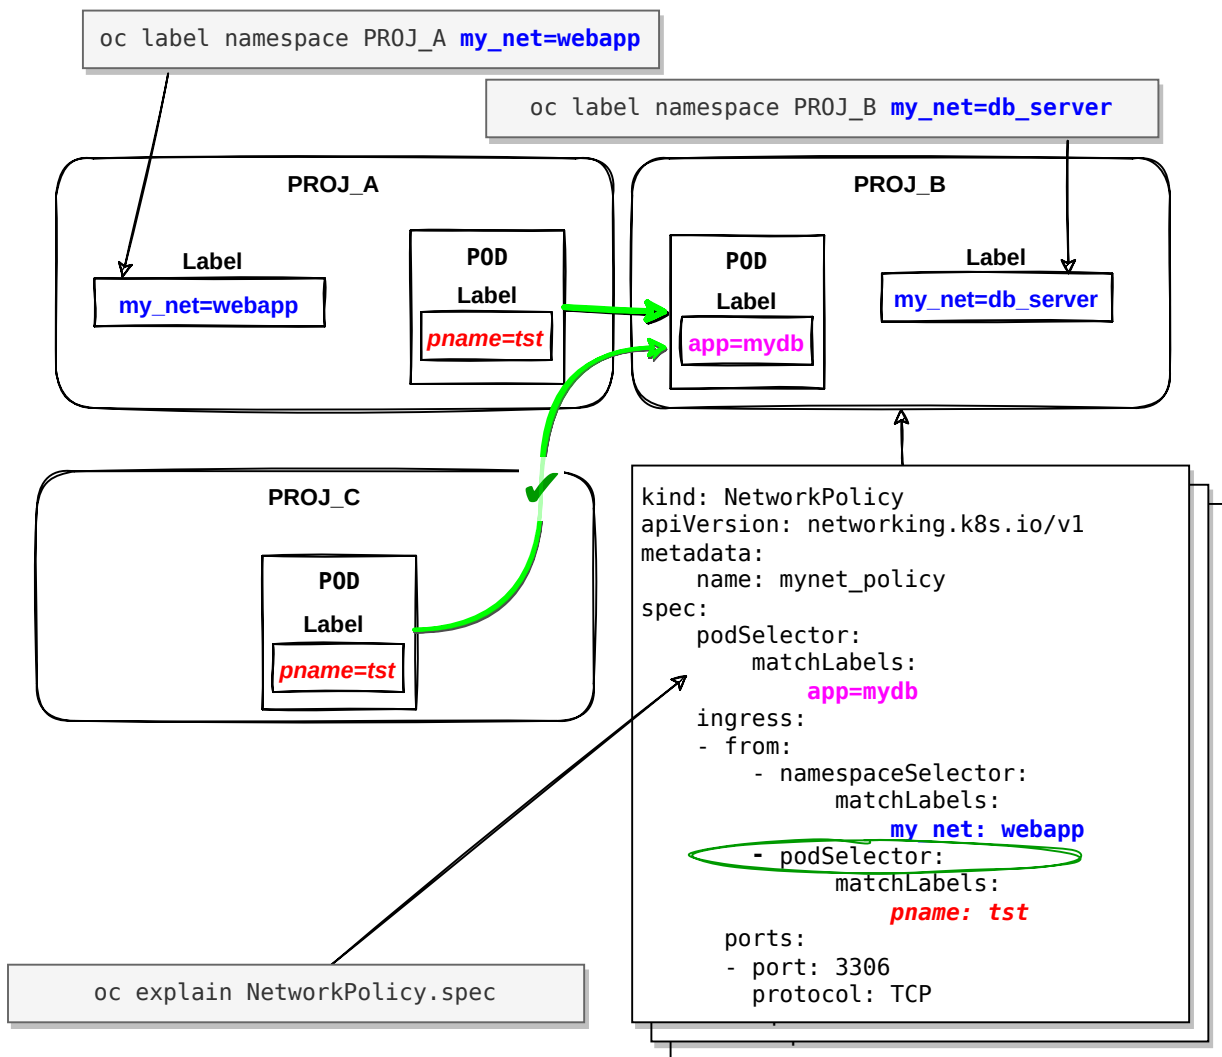

An operator is a control loop program and it can respond to events.  
It communicates with the API server to manage k8s resources





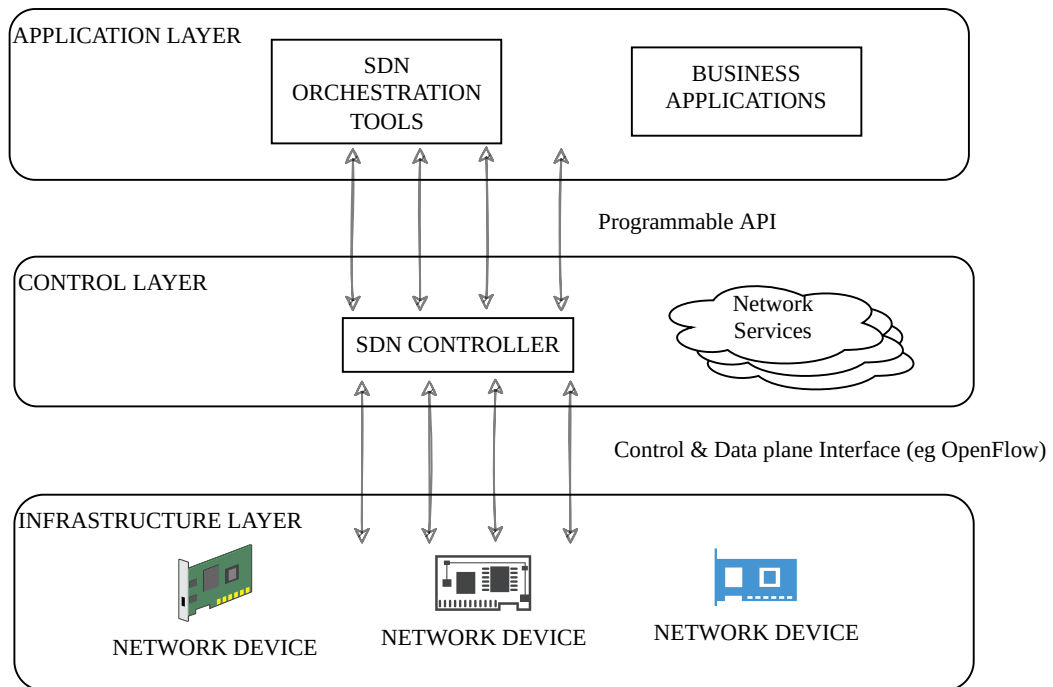




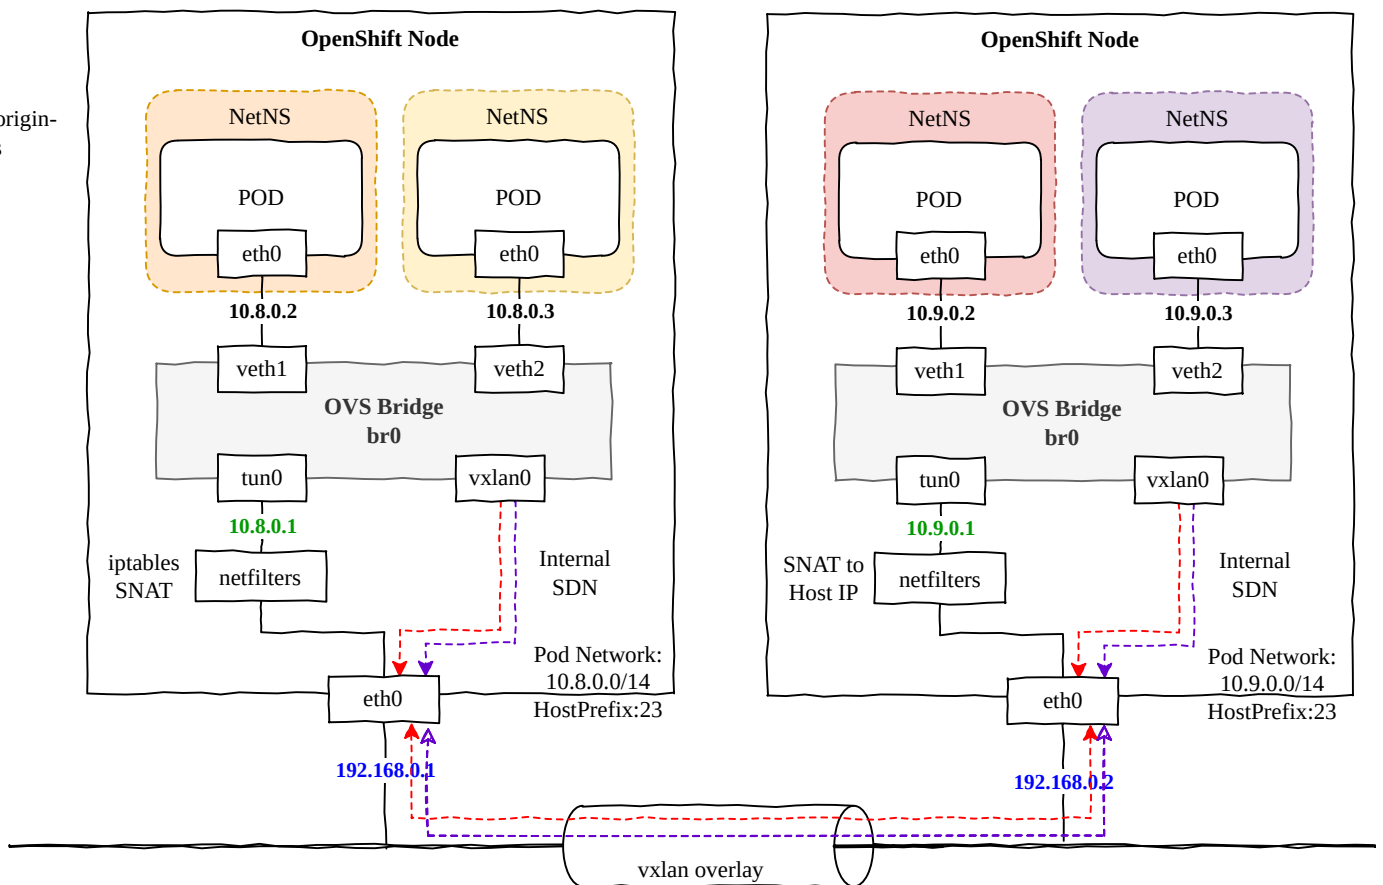


## SDN

- Abstraction of network layers
- Decouple network control and forwarding functions




use  
openshift/origin-  
tools



nnnnnnnn . nnnnnnxx . xxxxxxxx . xxxxxxxx

# POD Scheduling

1. Get a list of all NODES
2. Go through all the predicates for **FILTERing**. If NODE fails predicate rule, remove from list. Region affinity.
3. With remainder list of NODES, **prioritize** them using the weightage rules. NO filtering of NODES done here. Zone anti-affinity.
4. Select the NODE with highest points.



```
oc label node <NODE> <KEY>=<VALUE>
```

Region

<KEY> = failure-domain.beta.kubernetes.io/region

A set of hosts in closed geographical area. High speed connectivity.

Zone (availability zone)

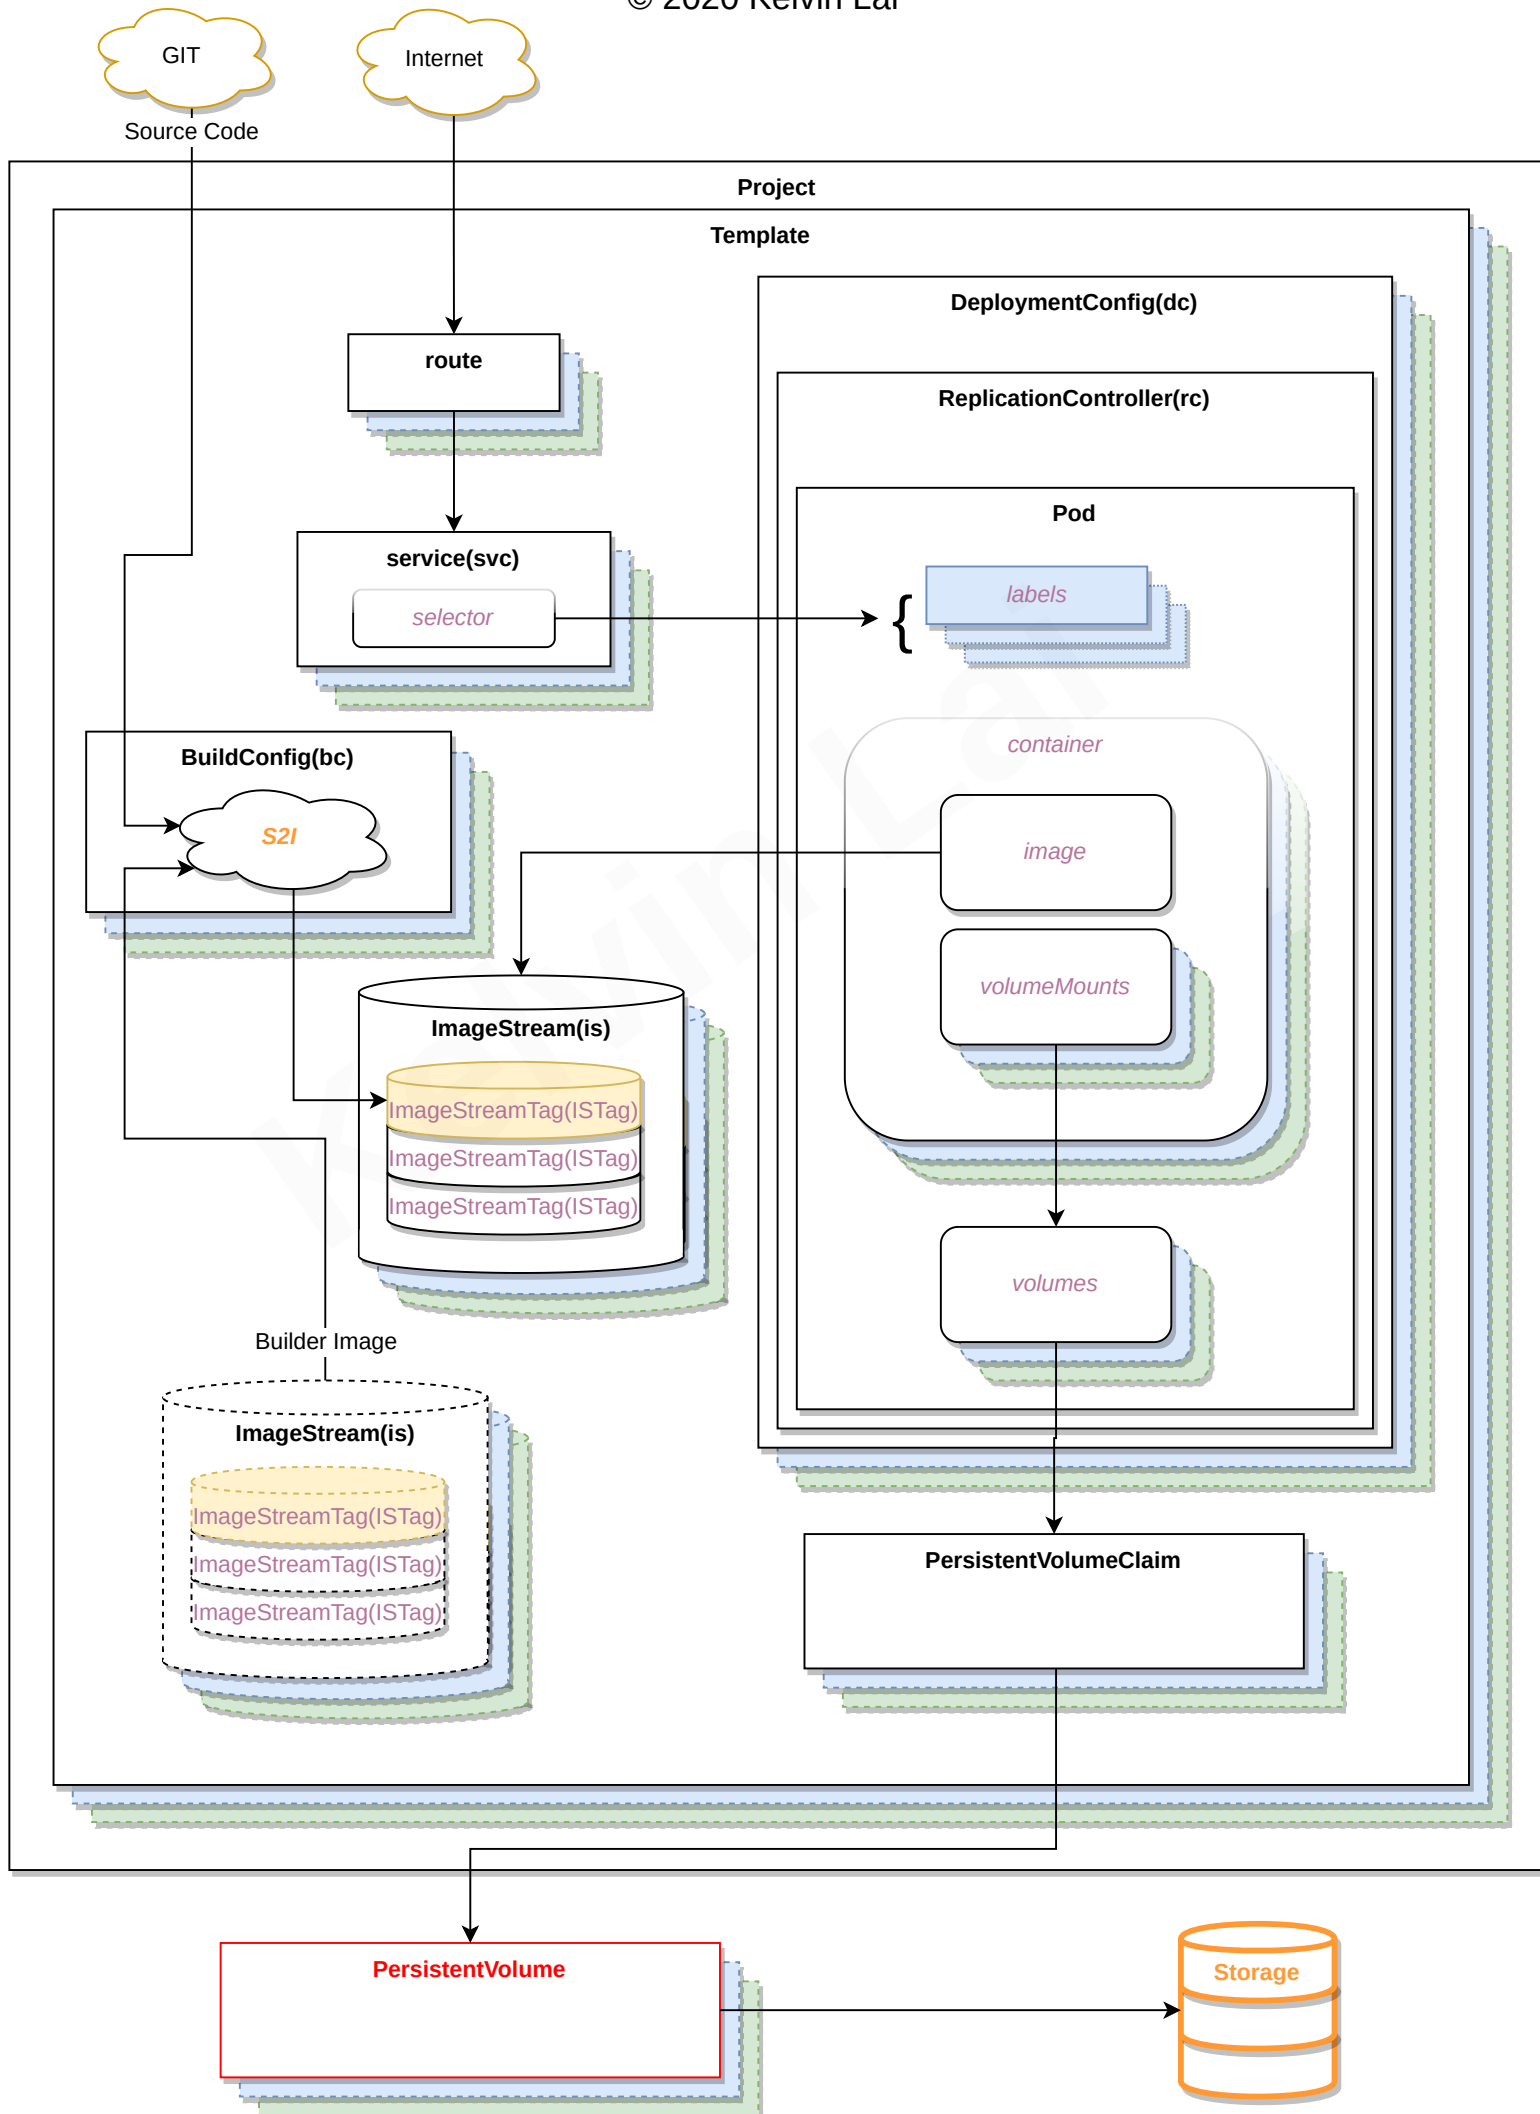
<KEY> = failure-domain.beta.kubernetes.io/zone

A set of hosts that share common critical infra components (ups, switch, storage)

**Upgrade Path Graph:** [https://access.redhat.com/labs/ocpupgradegraph/update\\_channel](https://access.redhat.com/labs/ocpupgradegraph/update_channel)

# OpenShift Resource Types

© 2020 Kelvin Lai



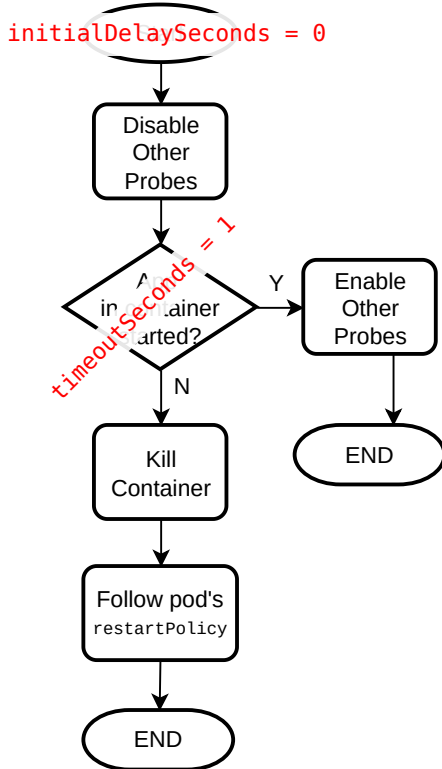
# Health Monitoring

## Probes

`oc set probe --help`

### Startup

Application in container started?  
spec.containers.startupprobe




*successThreshold = 1  
failureThreshold = 3*

**Mandatory**

**Optional**

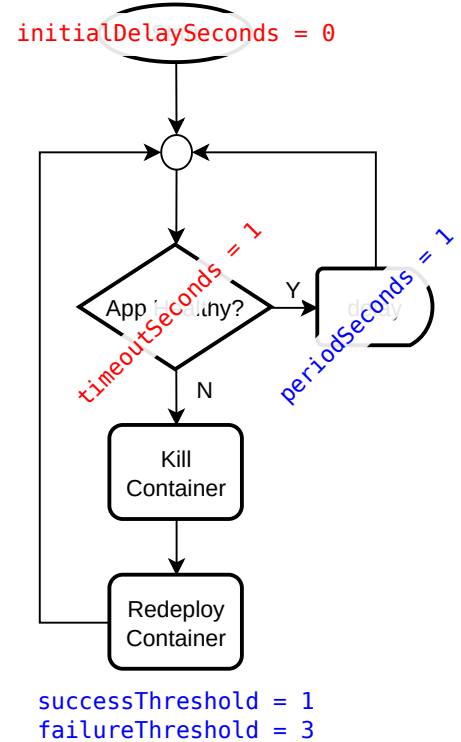
### Readiness

Container ready to serve requests?  
spec.containers.readinessprobe



### Liveness

Container application healthy?  
spec.containers.livenessprobe



## Deployment/DC

```

spec:
  template:
    spec:
      containers:
        - image: ...
          ...Probe:
  
```

```

readinessProbe:
  httpGet:
    path: /ready
    port: 8080
  initialDelaySeconds: 15
  timeoutSeconds: 1
  
```

```

livenessProbe:
  exec:
    command:
      - cat
      - /tmp/health
  initialDelaySeconds: 15
  timeoutSeconds: 1
  
```

```

livenessProbe:
  tcpSocket:
    port: 3306
  periodSeconds: 15
  timeoutSeconds: 1
  
```

```

oc set probe deployment test-php \
  --readiness \
  --get-url=http://:8080/ready \
  --initial-delay-seconds=15 \
  --timeout-seconds=1
  
```

```

oc set probe deployment test-php \
  --liveness \
  -- cat /tmp/health \
  --initial-delay-seconds=15 \
  --timeout-seconds=1
  
```

```

oc set probe deployment test-php \
  --liveness \
  --open-tcp 3306 \
  --period-seconds 15 \
  --timeout-seconds 1
  
```

## Blue-Green Deployment

### Deploy production app, green

```
oc new-app registry.example.com/myapp:v1 --name green
oc expose deployment green      # expose dc/deployment to get svc
oc expose svc green            # expose svc to get route
oc get route                   # get hostname to access app
```

### Deploy updated version of app, blue

```
oc new-app registry.example.com/myapp:v2 --name blue
oc expose deployment blue      # expose dc/deployment to get svc
```

```
oc patch route green -p '{"spec":{"to":{"name":"blue"}}}' # update route to use blue svc
```

## A/B Deployment

### Deploy production app, app-a

```
oc new-app registry.example.com/myapp:v1 --name app-a
oc expose deployment app-a      # expose dc/deployment to get svc
oc expose svc app-a --name myroute # expose svc to get route
oc get route                   # get hostname to access app
```

### Deploy updated version of app, app-b

```
oc new-app registry.example.com/myapp:v2 --name app-b
oc expose deployment app-b      # expose dc/deployment to get svc
```

Note: router balances traffic according to weights(default=1).

### Add service app-b to route and set weightage as 80% for app-a and 20% for app-b

```
oc set route-backends myroute app-a=80 app-b=20
oc set route-backends myroute # verify configuration
```

N-1 Compatibility & Graceful Termination - Refer to Notes

find out haproxy  
weight pct vs conn

# install-config.yaml

## Deployment Configuration

**Cluster Name**  
aio

**Base Domain**  
example.com

**Pull Secret**  
from console.redhat.com

**SSH Key**  
your public key

## Workload

**Node Config**  
control plane & compute

## Infra

**Platform**  
infrastructure type

**Platform Info**  
username/password

openshift-install

## Bootstrap Ignition


k8s manifest

## Control Plane Ignition

k8s manifest

## Compute Ignition

k8s manifest



#### Red Hat OpenShift Container Platform


- Public/private DC.
- Bare metal and multiple cloud and virtualization providers.
- Full control by customer.

#### Red Hat OpenShift Dedicated

- Managed cluster in public cloud.
- RH manages the cluster.
- Customer manages updates and add-on services.

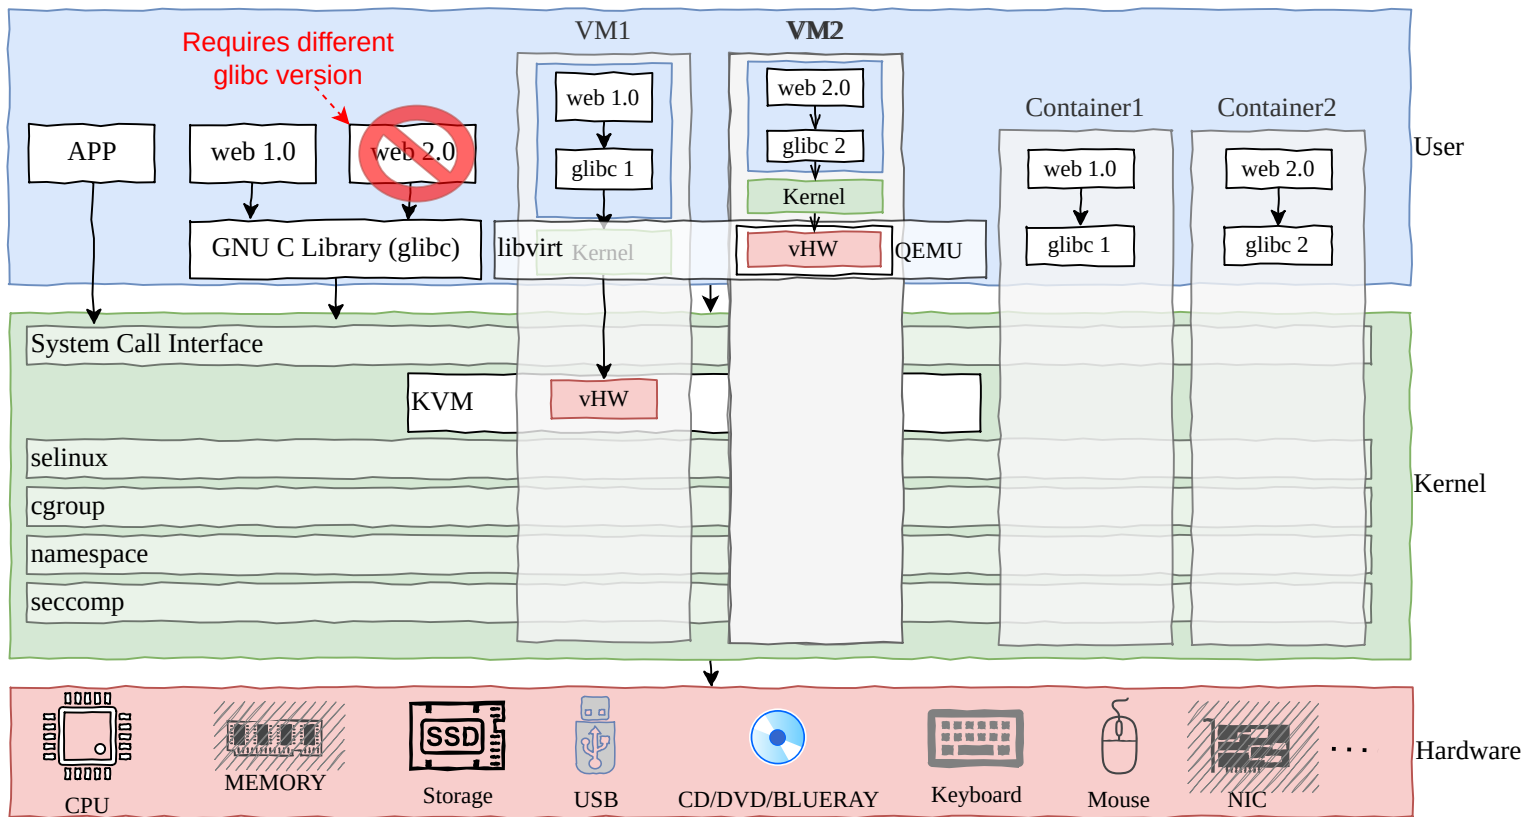
#### Red Hat OpenShift Online

- Public hosted cluster.
- Shared resources by multiple customers.
- RH manages cluster life cycle.

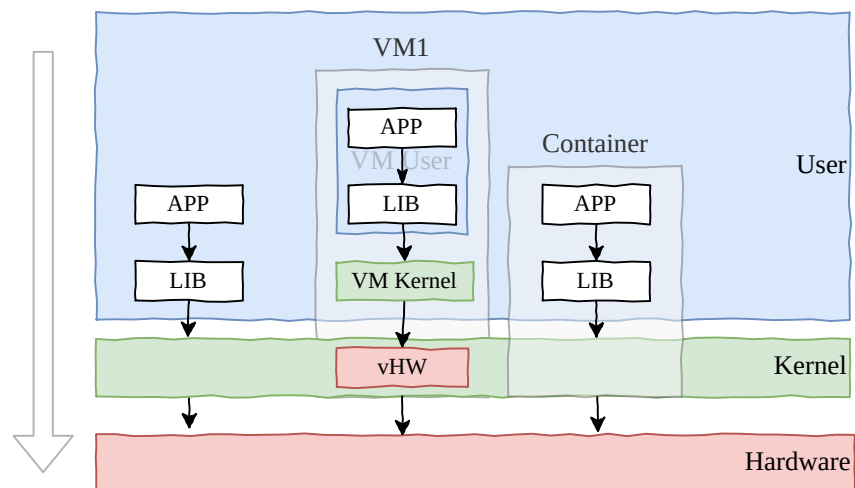




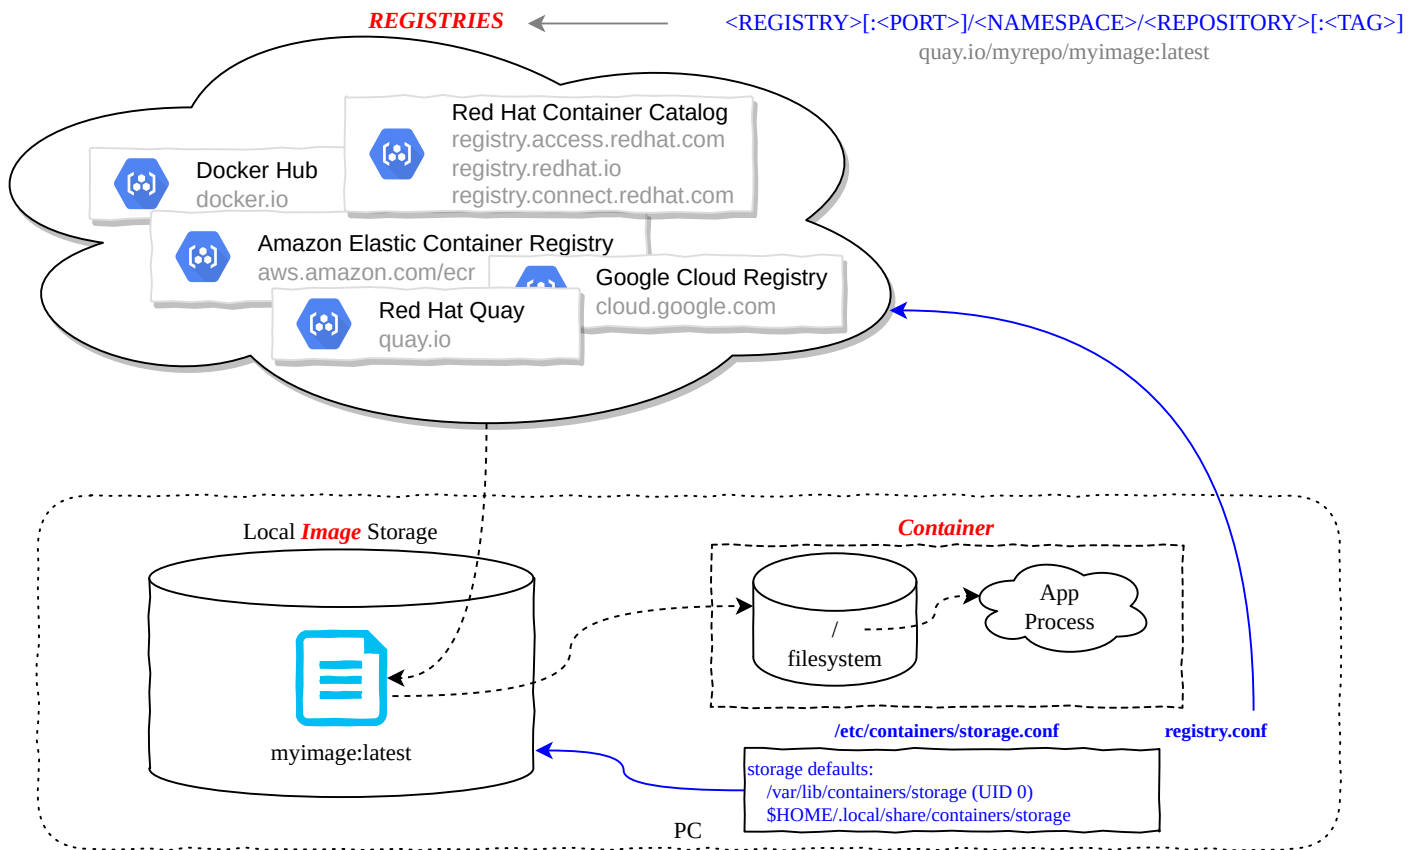
## VM vs Container



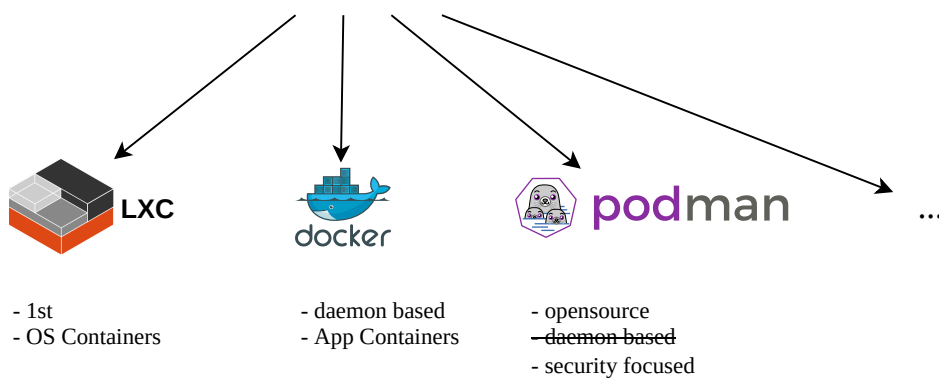
Ref: <https://www.redhat.com/en/blog/all-you-need-know-about-kvm-userspace>  
<https://www.packetcoders.io/what-is-the-difference-between-qemu-and-kvm/>



## Container Architecture



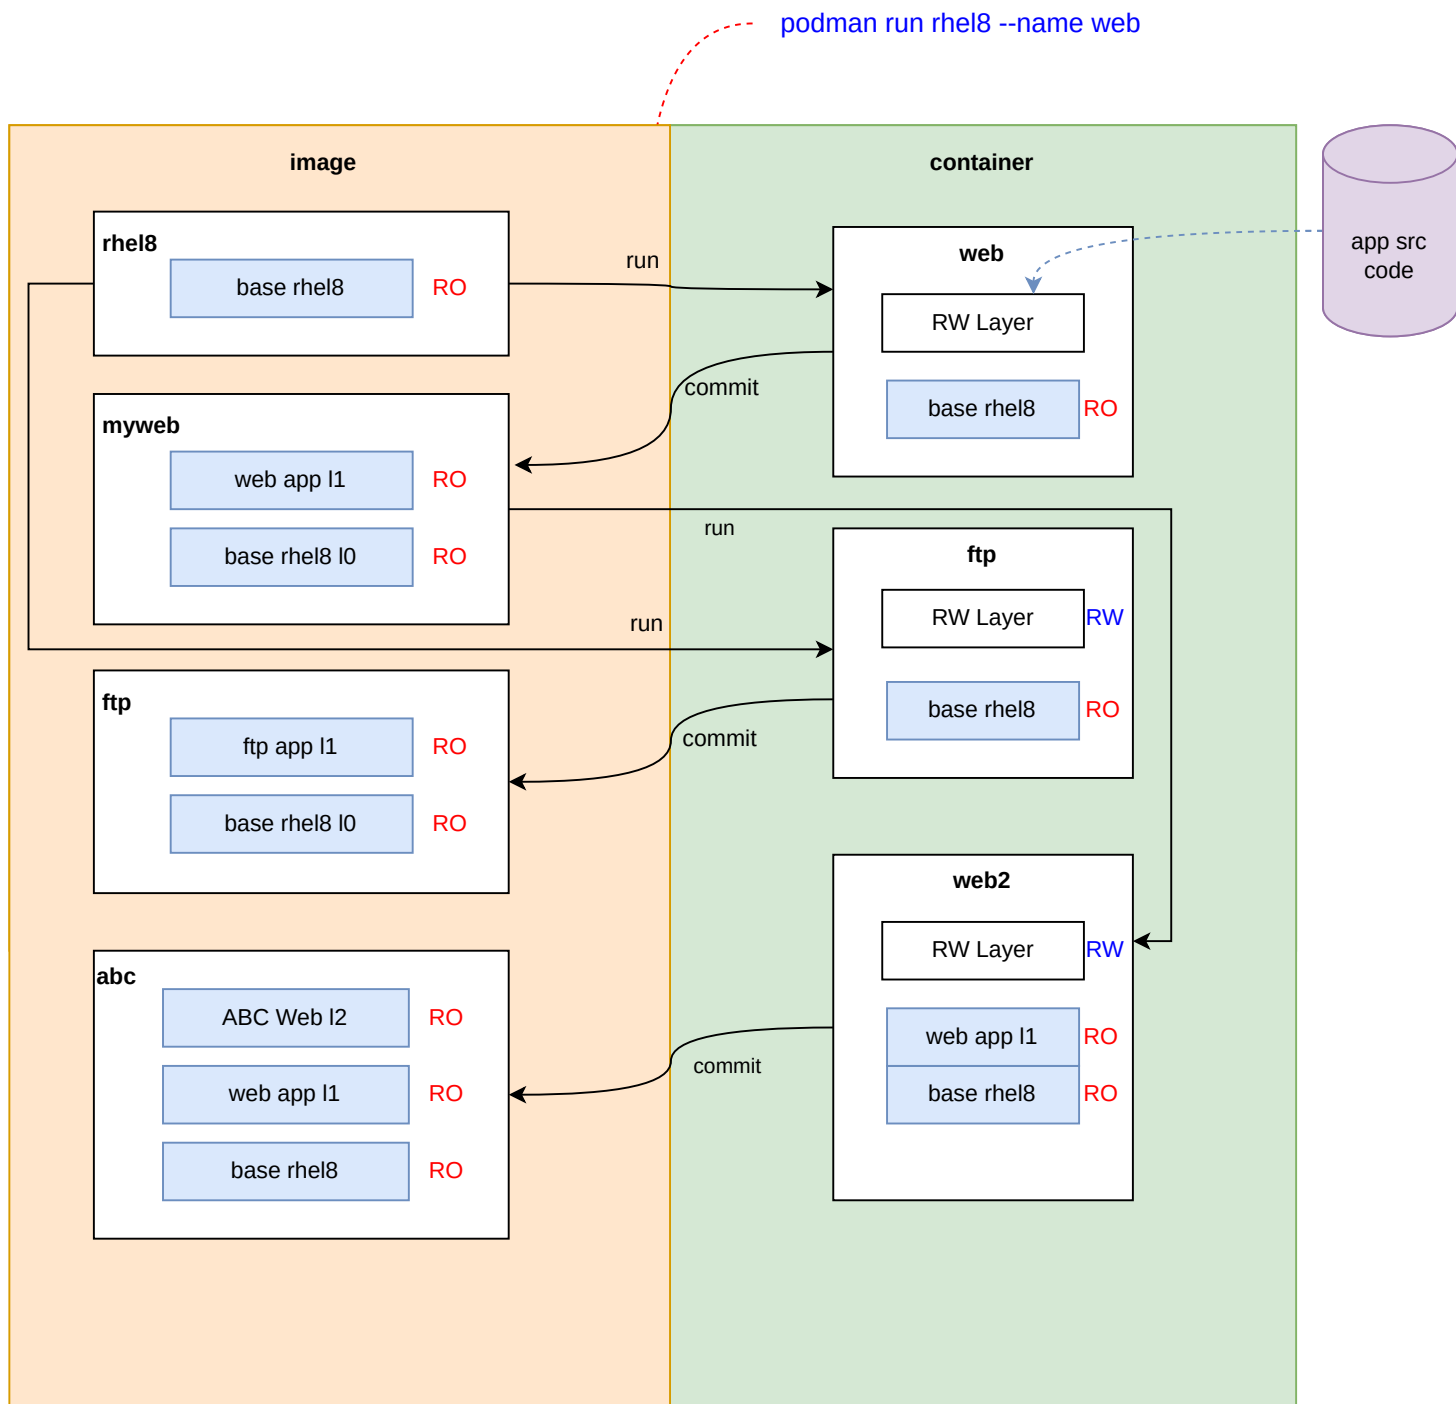
## Container Utilities



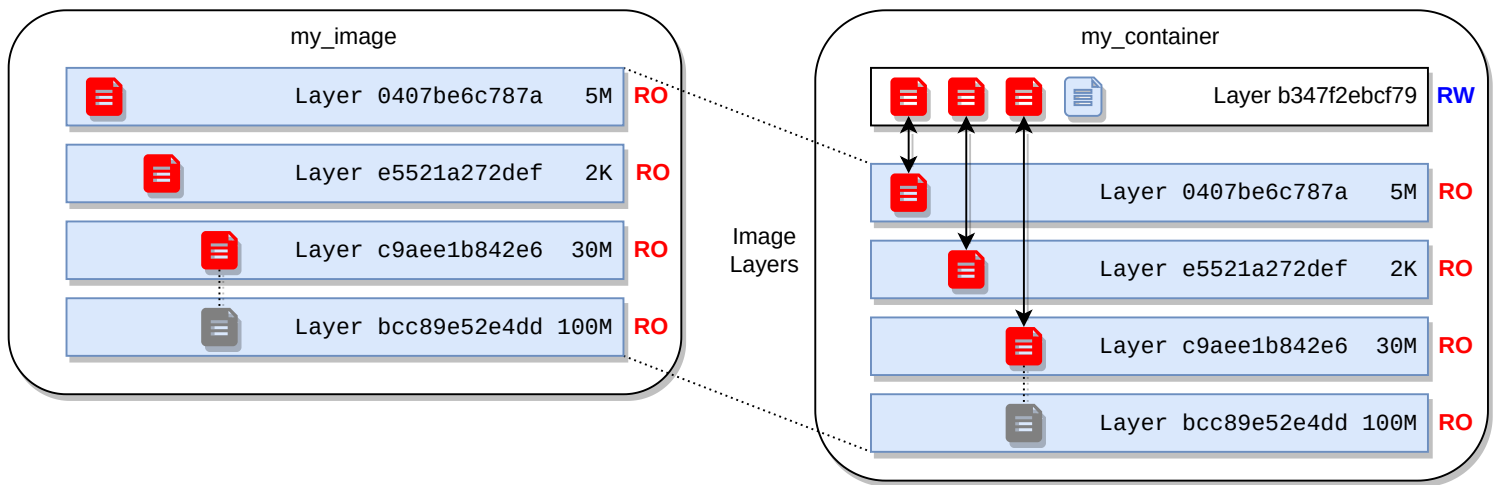
## OS Container Vs Application Containers

# Creating Image

1. Manual
2. Dockerfile/Containerfile
3. Source-To-Image(s2i/STI)
  - a) get runtime image and create container
  - b) clone source code into container
  - c) compile source code
  - d) deploy/publish compiled app
  - e) cleanup
  - f) save container as image

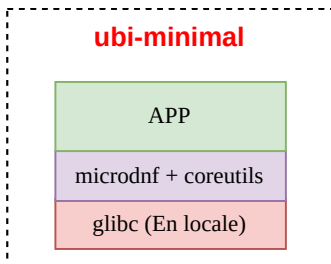


# UnionFS - A Stackable Unification File System



## BASE IMAGE TYPES


### MINIMAL



Designed for apps that contain their own dependencies (Python, Node.js, .NET, etc.)

- Minimized pre-installed content set
- no suid binaries
- minimal pkg mgr (install, update & remove)


### PLATFORM



For any apps that runs on RHEL

- Unified, OpenSSL crypto stack
- Full YUM stack
- Includes useful basic OS tools (tar, gzip, vi, etc)

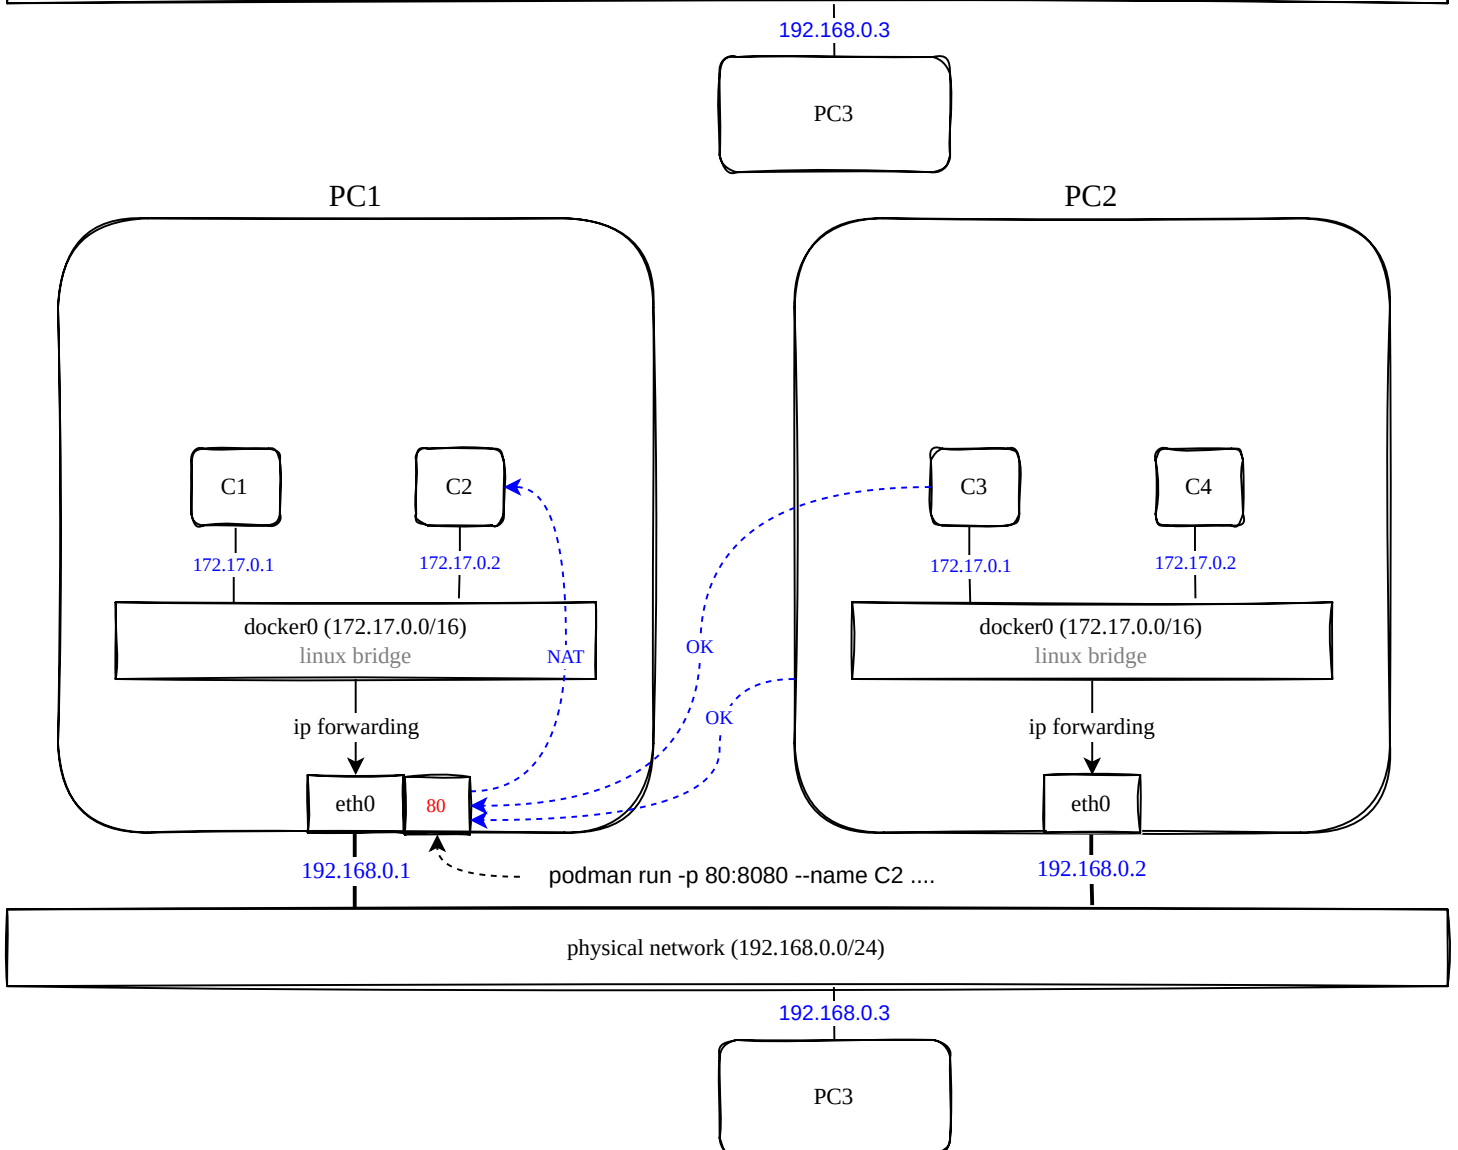
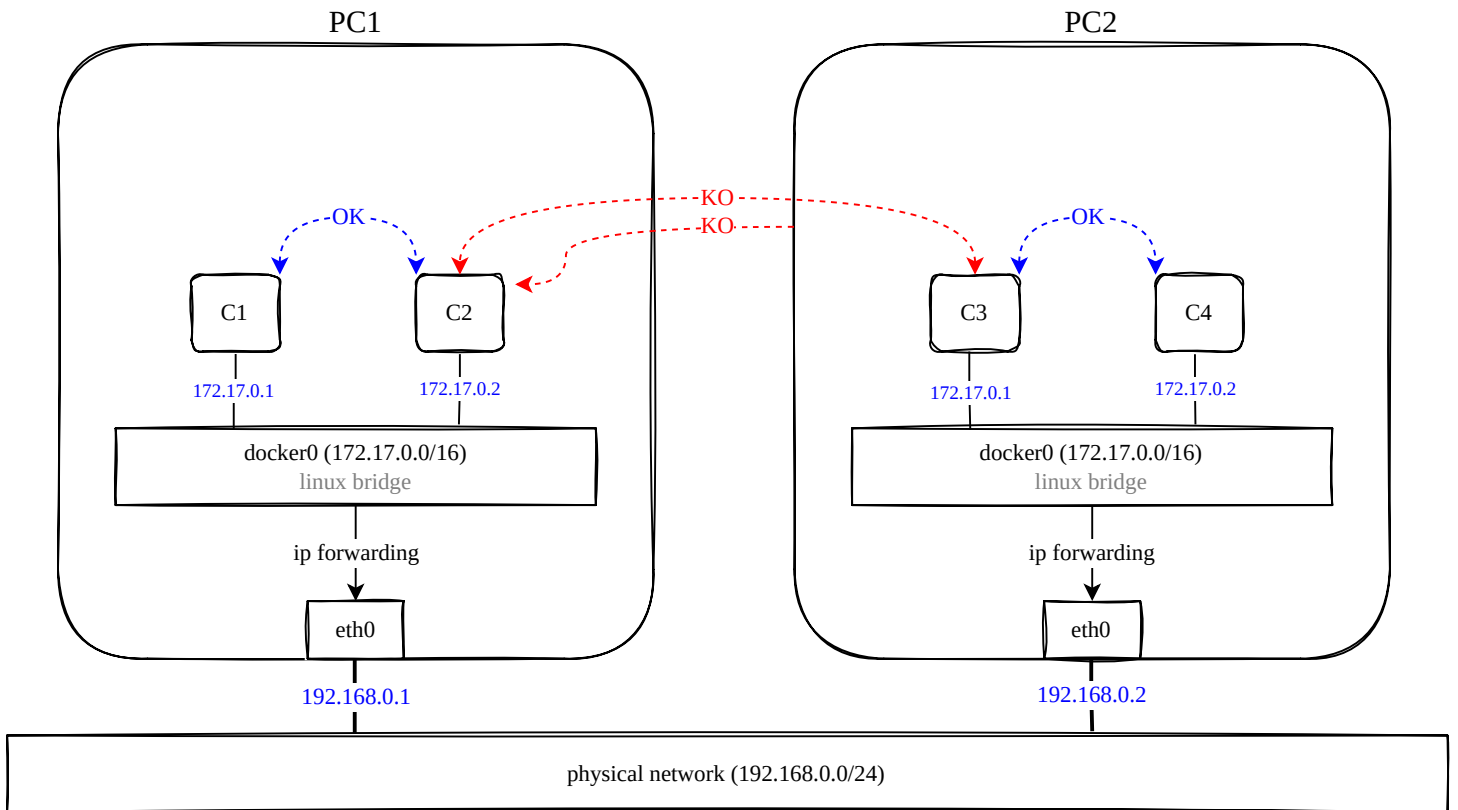
### MULTI-SERVICE



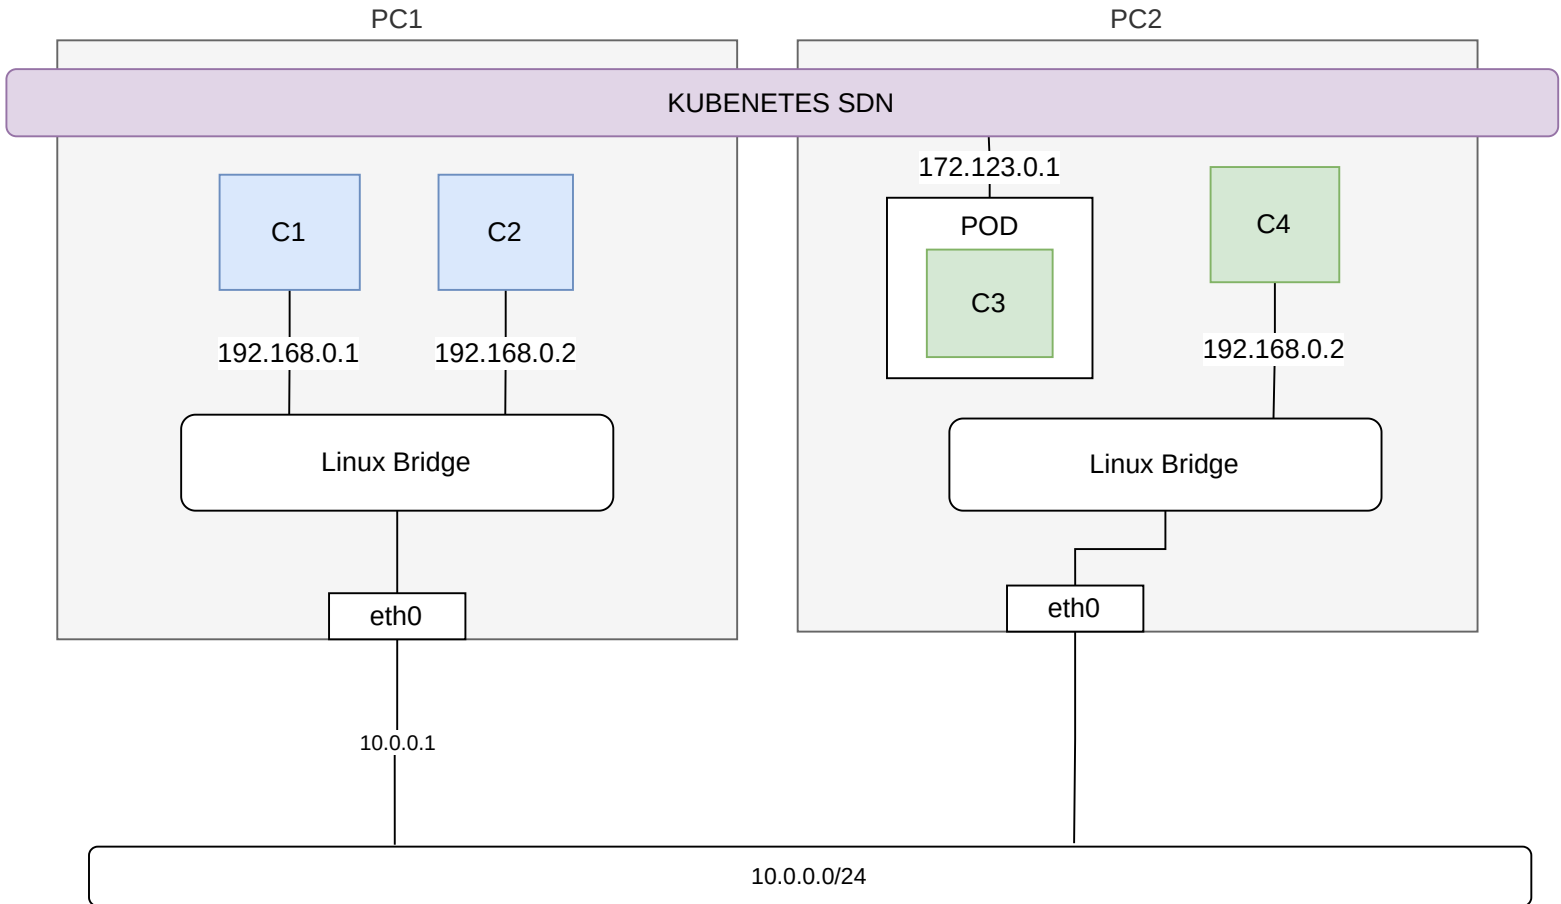
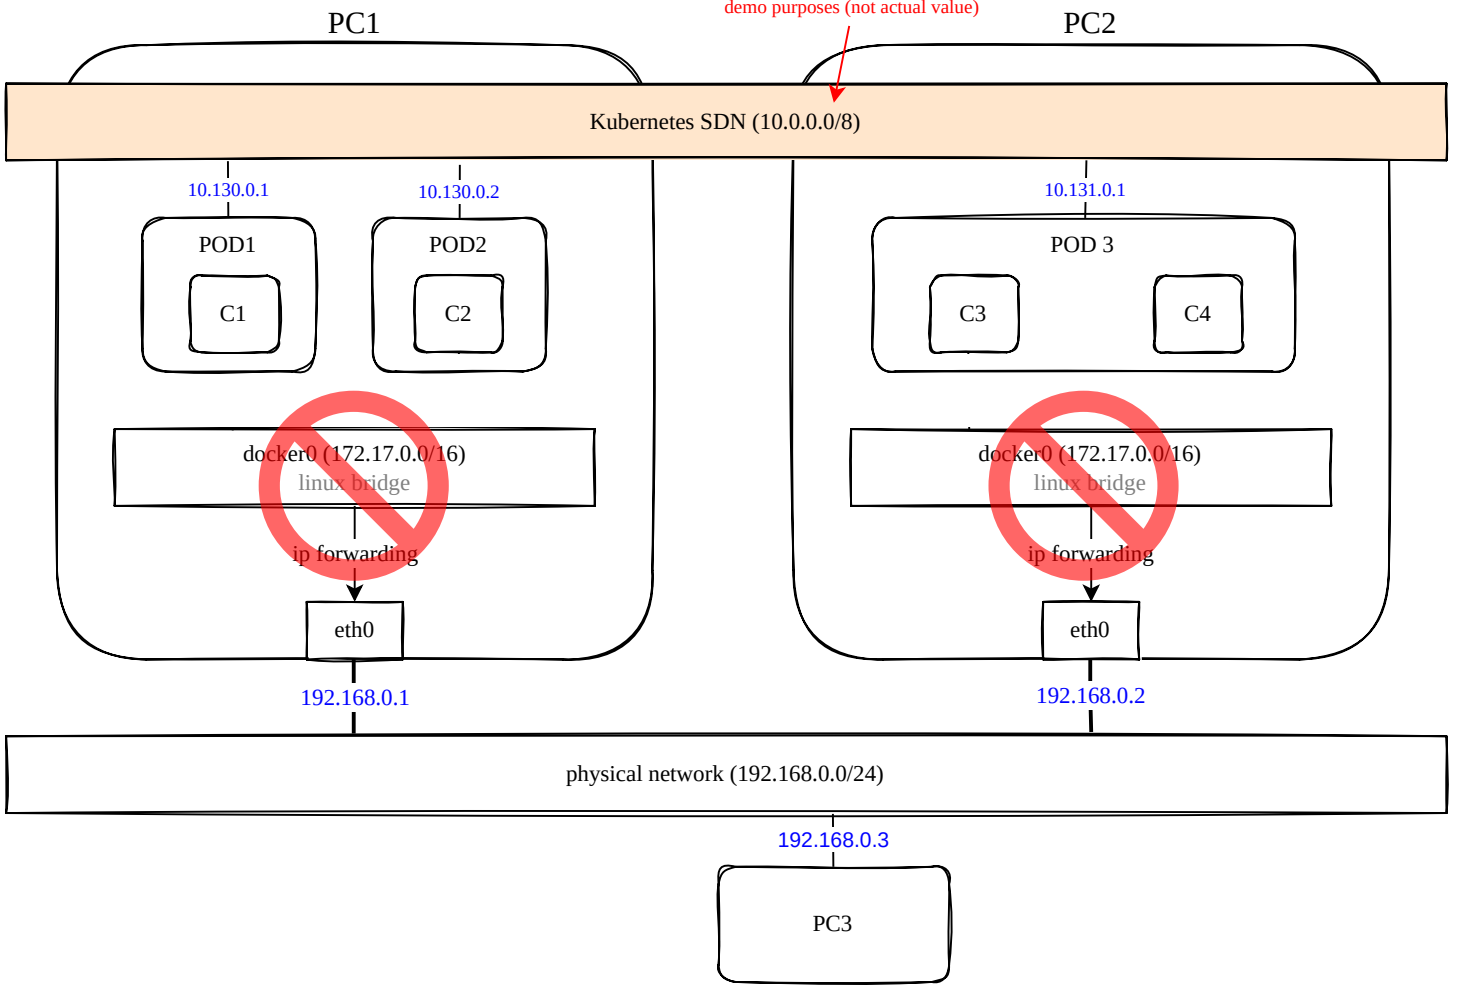
Eases running multi-service in single container

- configured to run systemd on start
- allows you to enable th services at build time

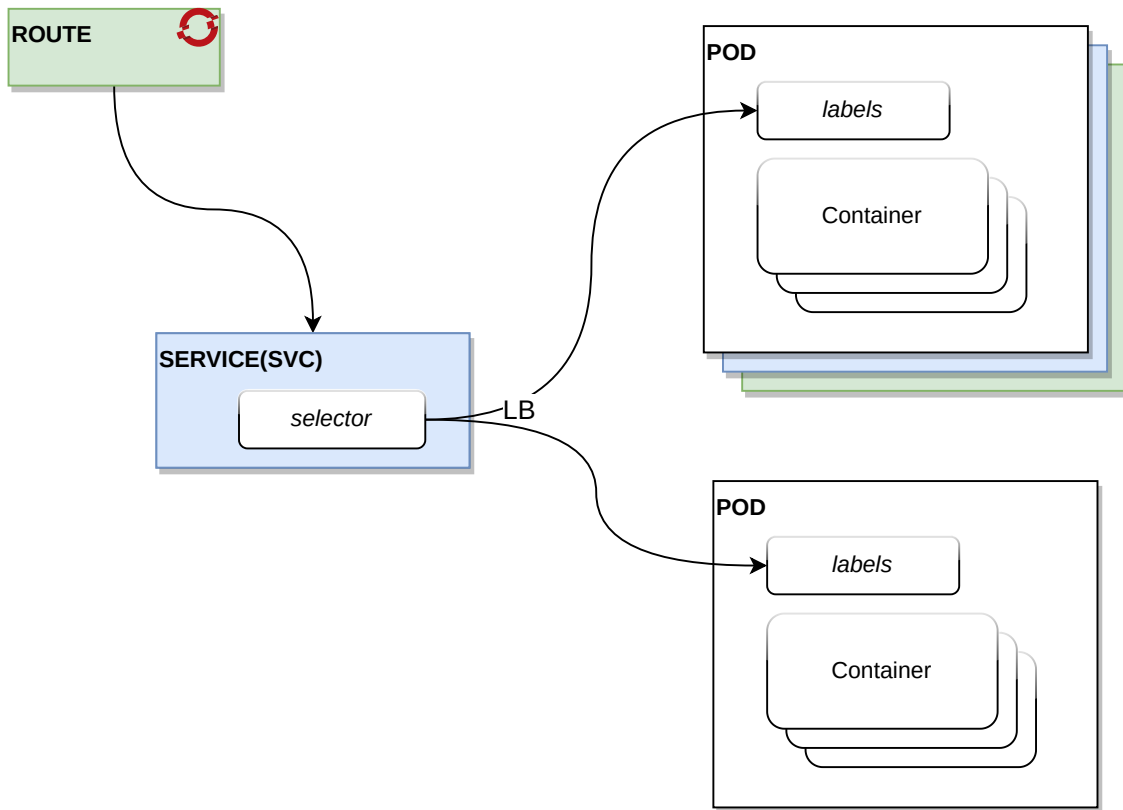
## Basic Network - Container vs Kubernetes



demo purposes (not actual value)



# Route, Service and Pod Relationship



## POD

A pod contains one or more containers.

## SERVICE

A service references the pod(s) by using the label selector.

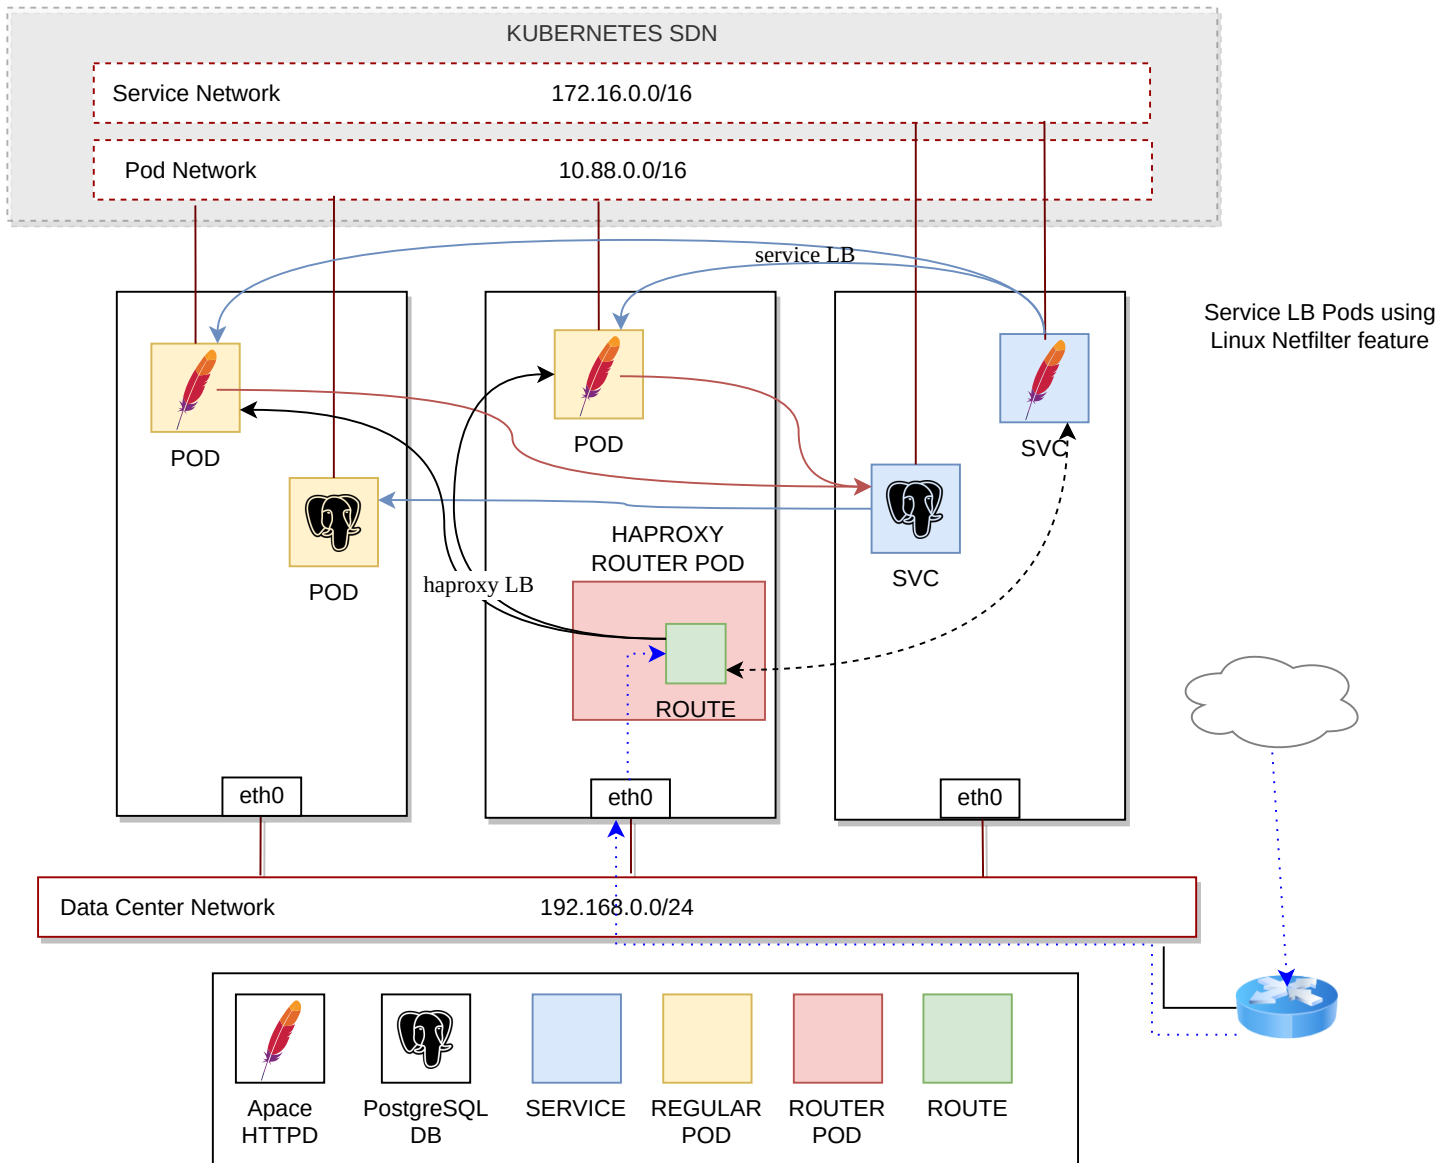
The service load balances the connections between all the pods.

## ROUTE

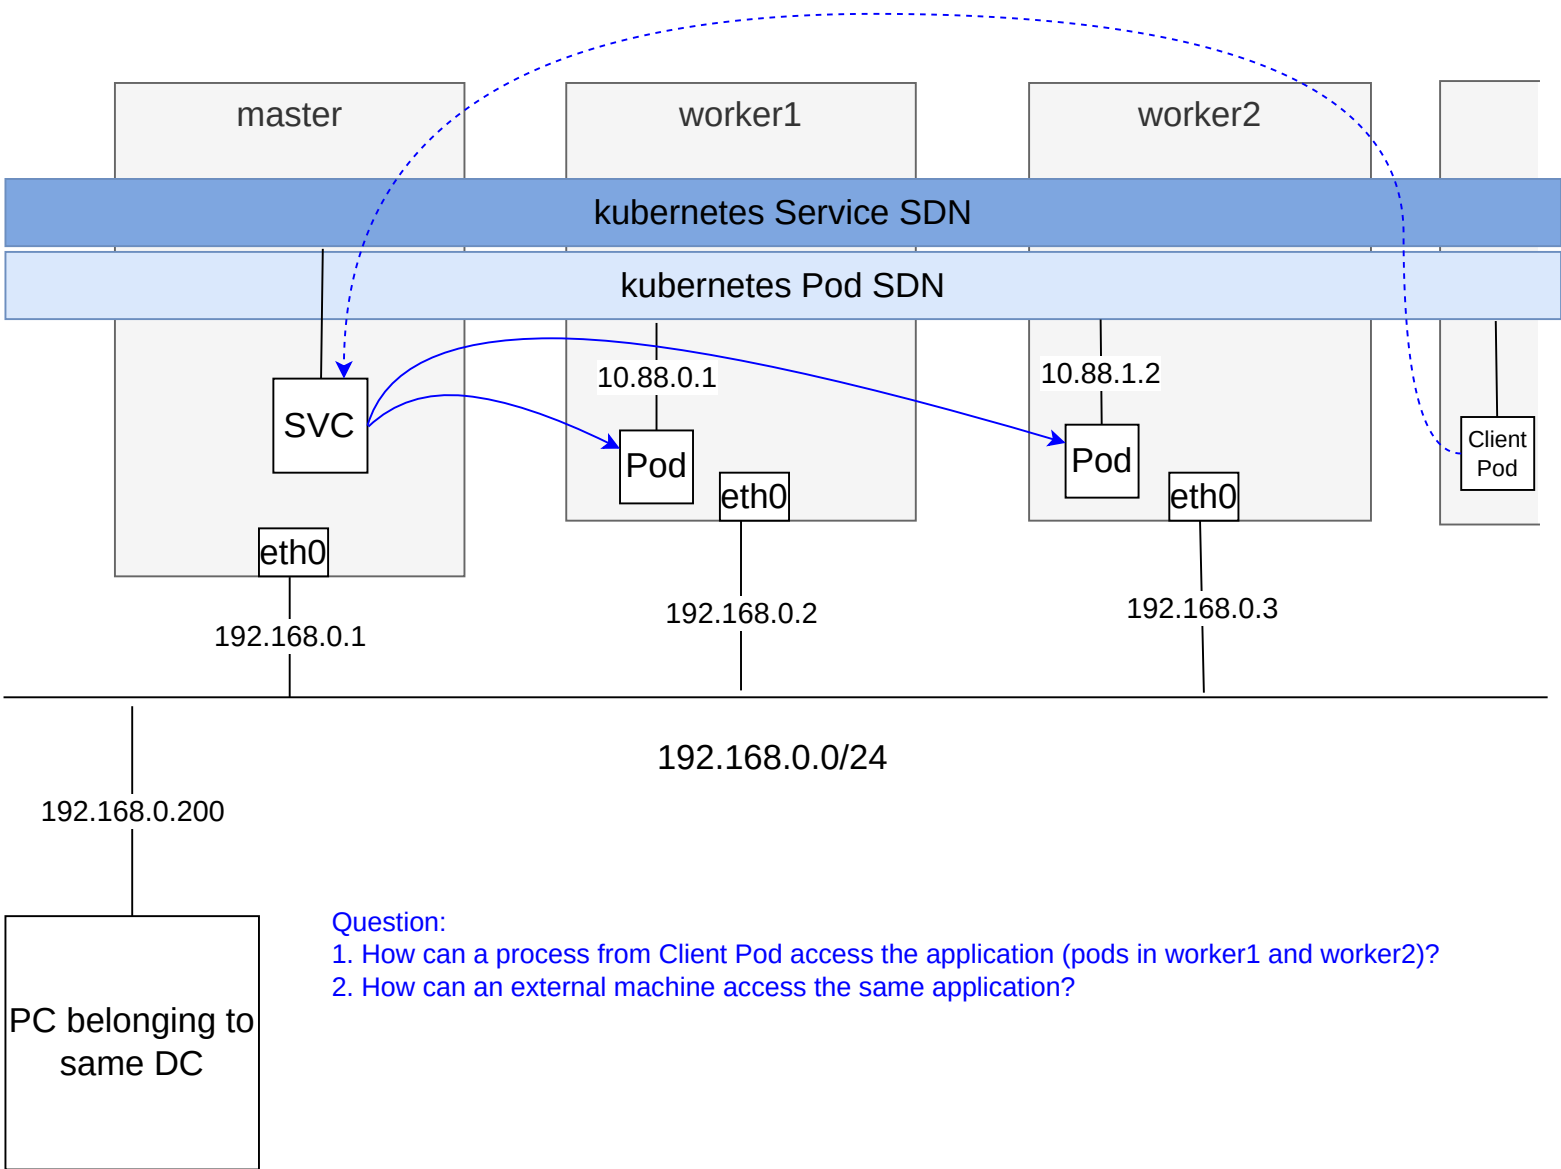
A route exposes the service to the external world.

**Warning:** A service "can" refer to different pods, if the pods have the same label.

## Sample of how Services are used



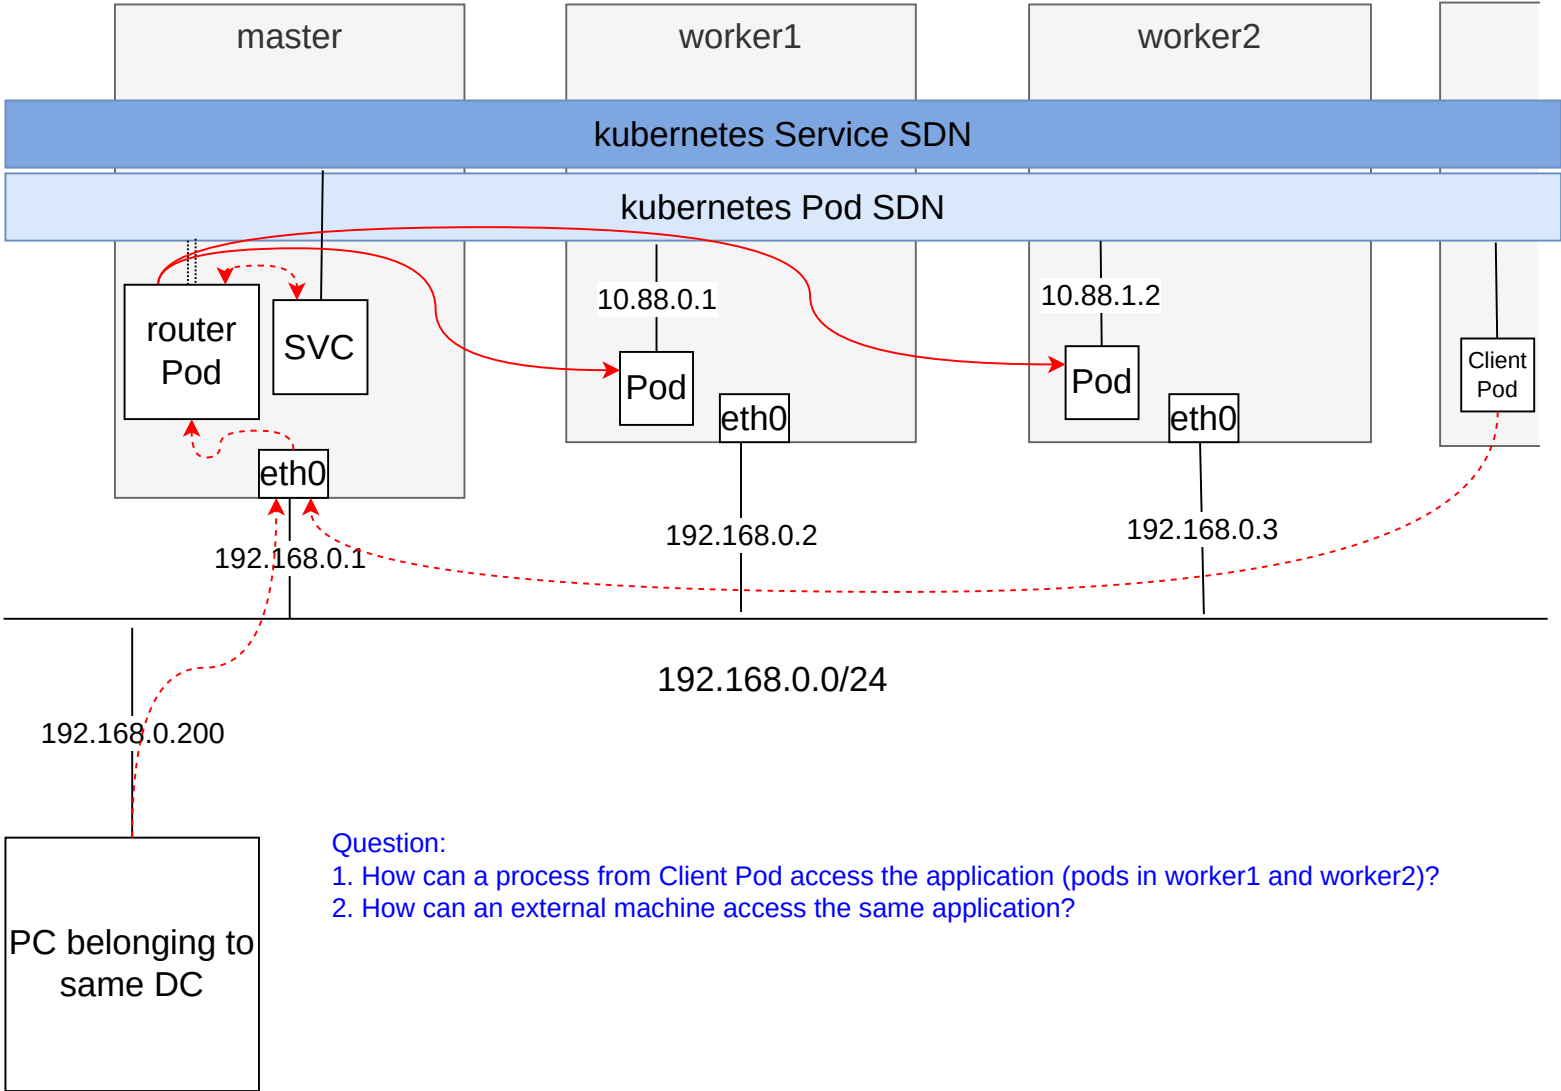




Question:

1. How can a process from Client Pod access the application (pods in worker1 and worker2)?
2. How can an external machine access the same application?





Question:

1. How can a process from Client Pod access the application (pods in worker1 and worker2)?
2. How can an external machine access the same application?

# Deploying Applications with OpenShift

Methods to create applications:

1. Using existing containerised applications

```
oc new-app --docker-image=<IMAGE>
```

2. From Source Code using S2I

```
oc new-app <URL>
```

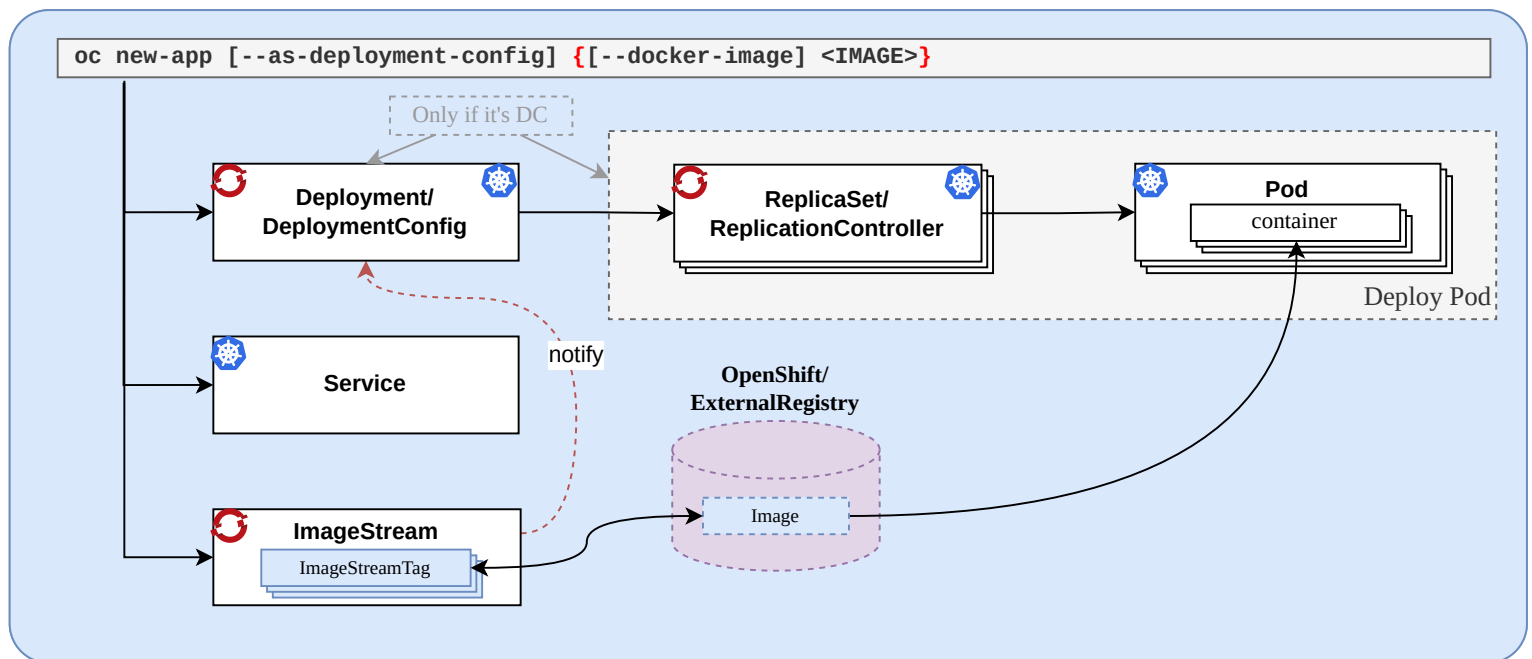
3. Using yaml/json file

```
oc new-app -f <FILE>.yaml
```

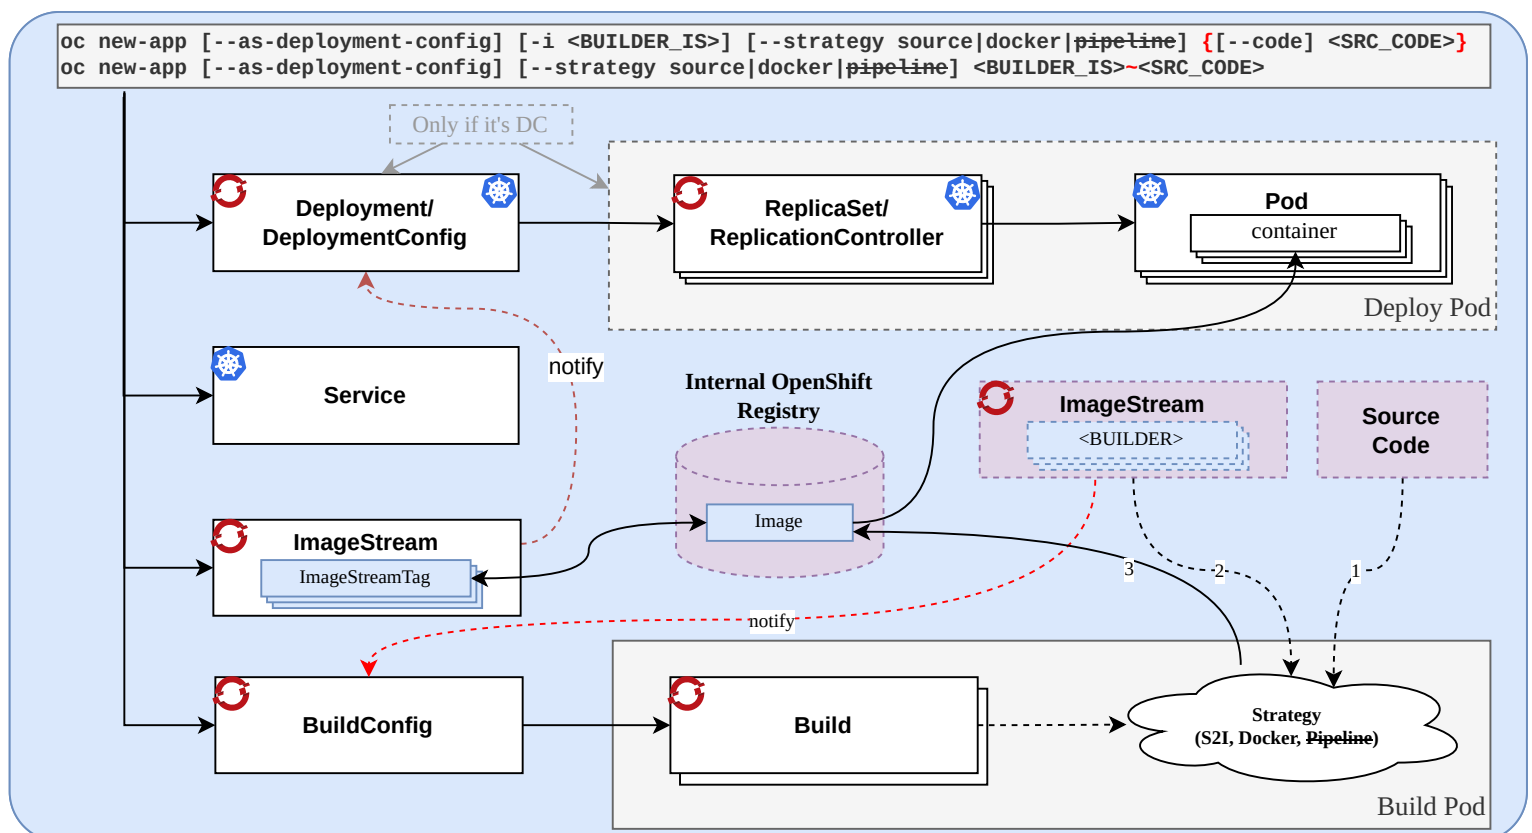
4. Using template

```
oc new-app --template=<TEMPLATE> --param=<PARAM> --param-file=<PARAM_FILE>
```

## 1. Use Existing Image



## 2. Managed Life Cycle



```
oc new-app -i myphp https://github.com/user/myapp#branch --context-dir <DIR>
```

```
oc new-app -i myphp:7.1 https://github.com/user/myapp
```

```
oc new-app myphp:7.1~https://github.com/user/myapp
```

NOTE: -i option needs git client to be installed

## Options

- o json|yaml inspect resource definitions without creating
- name <NAME> adds a label "app=<NAME>" to all resources, Use `oc delete all -l "app=<NAME>"` to cleanup

## IMPORT IMAGES

```
oc import-image <IMG_STREAM> [--confirm] --from <IMAGE> [--insecure]
where,
    <IMAGE> = <REGISTRY>[:<PORT>]/<NAMESPACE>/<REPOSITORY>[:<TAG>]
```

oc new-app command in OpenShift 4.5 makes use of [deployment](#) resource. Use --as-deployment-config if you wish to create [deployment config](#) instead.

SERVICE(SVC)

```
oc expose <DC/DEPLOYMENT/RC/RS/POD> <RESOURCE_NAME>
```

```
DNS NAME = <SVC>.<PROJ>[.svc.cluster.local]
ENVIRONMENT VARIABLE IN POD = <SVC>_SERVICE_HOST
```



ROUTE

```
oc expose svc <SVC_NAME> [--name <ROUTE_NAME>] [--hostname <FQDN>]
```

```
DNS DEFAULT NAME = <ROUTE_NAME>.<PROJ>.<DOMAIN WILDCARD>
<DOMAIN_WILDCARD> = apps.<BASE_DOMAIN>
```

© 2020 Kelvin Lai



Container Image Naming Convention: <REGISTRY>[:<PORT>]/<NAMESPACE>/<REPOSITORY>[:<TAG>]

podman pull [quay.io/myuser/mysql-57-rhel7](#)

